

Good News and Bad News: The Effects of Consuming Professional News Outlets During the 2022 Brazilian Elections*

Caio Castro[†] and João F. Pugliese[‡]

Abstract

We study the effects of direct consumption of professional news in a randomized experiment with 1,658 voters during Brazil’s 2022 elections. Participants interested in subscriptions chose among five newspapers. Offering three-month free subscriptions and reading incentives shifted consumption from social media to newspapers without increasing total news consumption. Treatment improved news knowledge and trust in professional media, with evidence of stronger support for democratic norms and no effect on affective polarization. Effects on trust and democratic attitudes were concentrated among non-left-wing participants. Despite the positive overall effects of consumption, only three participants subscribed to newspapers after the intervention.

*We thank Matthew Gentzkow for his guidance and support throughout this project. The project benefited from conversations with Arun Chandrasekhar, Ernesto Dal Bó, and Ro’ee Levy. We thank Leon Eliezer, Luigi Garzon, and Bernardo Jazra for exceptional research assistance. We thank seminar audiences at DevPEC and PacDev for helpful comments and feedback. This project was made possible through the support of the Stanford King Center on Global Development’s Graduate Student Research Funding. The study was pre-registered in the AEA RCT registry (AEARCTR-0009940). The study protocol was approved by Stanford’s IRB panel (IRB-64785).

[†]São Paulo School of Economics, caio.castro@fgv.br

[‡]Stanford University, jfpugli@stanford.edu

1 Introduction

Only a minority of citizens consume professional news outlets directly. Most outlets sit behind paywalls, and less than 20% of adults in the United States – and only about 1 to 2% of adults in Brazil, the setting of this study – pay for a newspaper subscription.¹ Citizens increasingly get their news from other sources, such as semi-professional creators found on social media platforms. At the same time, trust in traditional news media has declined worldwide (Newman et al., 2022). Yet promoting professional news is suggested as a way to reduce misinformation and strengthen democracy (Council of Europe, 2022; Stiglitz et al., 2022).

Does directly consuming professional news improve citizens’ information and political attitudes when most news can also be obtained indirectly through social media and other secondary sources? The answer to this question is empirically unclear. Since information produced by professional outlets is not strictly excludable, it can spread across citizens even if paywall-protected. Indirect consumption through secondary sources can be a substitute for direct readership, preserving much of the information provided by newspapers (Galeotti and Goyal, 2010). Furthermore, the worldwide decline in readership could be consistent with consumers not obtaining much value from direct consumption.

This paper reports the results of a randomized controlled trial addressing this question, conducted during the 2022 Brazilian elections. We experimentally promote direct consumption of professional news by providing free subscriptions and small incentives. Participants choose which newspaper to receive, so the intervention subsidizes direct news consumption without assigning consumers to a particular outlet. Using pre-registered outcomes, we show that shifting news consumption from secondary sources toward professional newspapers improves news knowledge, increases trust in professional media, and, among centrist and right-wing participants, strengthens support for democratic norms, but generates virtually no lasting demand for paid subscriptions.

A challenge in studying professional news consumption outside of social media lies in the low take-up of professional news.² We make two choices in the study design to ensure naturalistic variation and high take-up. First, we recruited 1,658 participants on social media and screened for those who declared interest in a news subscription. This sample represents marginal consumers who are in principle willing to consume at least one of the professional outlets we partnered with, but lack a subscription. Second, rather than randomly assigning participants to a newspaper, at baseline we asked participants to choose their preferred option among five large professional outlets. Participants’ choices were strongly based on their perceived ideological congruence with the outlet, in line with existing evidence of demand for partisan news (e.g., Gentzkow and Shapiro, 2010; Chopra et al., 2024). We then randomly allocated participants into three groups: (i) control (55%); (ii) free subscription (5%); (iii) free

¹Brazilian data come from a 2022 report from Instituto Verificador de Comunicação, a non-governmental organization that measures media consumption among major news organizations. In the United States, data from the Alliance for Audited Media (AAM) compiled by the Pew Research Center show between 8 and 20% of adults possess a newspaper subscription, depending on the methodology.

²Batista Pereira et al. (2023) find an 11.4% redemption rate when offering a digital news subscription to a Brazilian outlet to a representative sample. Hopkins and Gorton (2024) find a 1.7% redemption rate when offering a local newspaper to a representative sample in Pennsylvania, United States.

subscription and incentives (40%), small monetary rewards for reading that nudge consumption. The design thus promoted professional news consumption aligned with participants' own choices and led to a large take-up of subscriptions (70.0%).

The 2022 Brazilian elections provide a particularly relevant setting to study the effects of direct professional news consumption. Brazilians are heavy social media users and barely pay for newspaper subscriptions, making it especially important to ask whether direct consumption of professional outlets matters.³ The election also took place in a highly polarized political environment and was followed by protests and roadblocks by supporters of the defeated candidate, culminating in the January 8, 2023, attacks on Congress, the Presidential Palace, and the Supreme Court.

We use administrative and survey data and a set of pre-registered outcomes to analyze the consequences of the treatment. We measure coupon uptake and outlet consumption using newspaper-provided administrative data. We combine a baseline and an endline survey conducted three months after the intervention (about a month after the elections) to measure outcomes. We also employ short WhatsApp surveys to measure news consumption behavior and news knowledge over the course of the intervention. Outcome families are summarized with variance-weighted indices (Anderson, 2008).

The treatment shifted news consumption from social media toward professional newspapers without increasing total news consumption. Newspaper administrative data further show that treated participants spent about 10 minutes per week on their chosen outlet's website or app. At endline, we estimate a precise zero effect on total news consumption (0.3 minutes per day, or 0.01 s.d.; $p = 0.88$), alongside a marginally significant decline in reported social media news consumption of 0.08 standard deviations ($p = 0.09$). Surveys conducted every two weeks show a similar pattern throughout the intervention. We therefore interpret the treatment as changing the sources from which participants obtained news, rather than increasing the total amount of news they consumed.

The incentives raised the share of participants who consumed any news from their outlet but not consumption conditional on reading. Below, we focus on results for the larger subscription-plus-incentives group rather than pooling the two arms.

Treatment participants scored 0.10 standard deviations higher on a news-knowledge index combining three news quizzes ($p = 0.06$).⁴ The quizzes were composed of true and false headlines, including actual fake news that was circulating.⁵ In the second quiz, conducted right before the second round of the elections, treated participants performed 0.13 standard deviations better ($p = 0.03$). At the question level, the treatment increases the probability of getting an answer correct by an average of 5.5 pp. ($p < 0.01$), driving higher scores, but also increases the probability of an incorrect answer by about 3 pp. ($p < 0.01$) driven by lower "I don't know", a sign of increased confidence in answering. We find null effects on general political knowledge,

³In 2022, about 87% of Brazilians aged 10 or older used the internet, and about 84% among those reported using social media (Instituto Brasileiro de Geografia e Estatística, 2023).

⁴Quizzes were time-limited, and performance was not rewarded monetarily, to discourage searching for answers.

⁵The fake news items were selected with assistance from journalists of *Estadão Verifica*, the fact-checking arm of one of the newspapers.

which included statements about the political and electoral system more broadly (e.g., how many senators are elected per state or who was the current president of the House or Senate).

The treatment strongly increased trust in both the chosen newspaper and the press. The treatment raises trust in the chosen outlet by 0.30 standard deviations ($p < 0.01$). The effect is particularly strong among those with low trust or who were unfamiliar with outlets (0.71 s.d.). An index summarizing favorable opinions of journalists and the press increases by 0.09 standard deviations ($p = 0.02$). Meanwhile, trust in news from social media and from friends and family moves in the opposite direction (-0.08 s.d., $p = 0.10$). The pattern indicates that participants updated on the relative quality of professional news. It also speaks to the worldwide decline in trust in media, which has coincided with declining exposure. While we can hypothesize that decreases in trust lead to a decrease in consumption, our experiment shows the relationship can also run in the opposite direction: increasing exposure can lead to an increase in trust.

Beyond outcomes directly related to the news or newspapers, the treatment also appears to strengthen support for democratic norms, although the evidence is less precise. To capture democratic norms, we measure voters' willingness to vote for hypothetical candidates who commit six undemocratic actions, following Graham and Svobik (2020). Summarizing the six items in an index, the treatment decreased willingness to vote by about 0.06 standard deviations on average, although the estimate is imprecise ($p = 0.13$). An index capturing trust in Brazil's electoral system was, on average, unchanged by the treatment. The clearest effect is on a specific proposal involving printing of voting receipts, put forward by the Bolsonaro government and regarded as promoting distrust of the country's electronic voting system. Support for the proposal fell by 0.09 standard deviations, about 4 percentage points ($p = 0.04$).

The treatment also increased political engagement without increasing affective polarization. Participants report following politics more closely (0.10 s.d., $p = 0.04$), and a pre-registered index of eight political activities increases by 0.06 standard deviations ($p = 0.11$). Election-specific activities, such as watching debates and persuading others to vote, generally move in the same direction. In contrast, we find precise null effects on affective polarization across measures involving the two presidential candidates and left- and right-wing politicians and supporters: the confidence intervals rule out changes larger than about 0.10 standard deviations.

Among the pre-registered moderators, we find that increases in trust in the press and elections, and the decrease in the willingness to vote for undemocratic politicians are concentrated in participants identifying as centrists or right-wing. This result suggests one possible mechanism where direct consumption of professional news, even if ideologically congruent, can balance out other politically slanted information received during the electoral period.

In spite of the strong increases in trust, there was scarcely any demand for subscriptions after the intervention. As subscriptions expired at the end of November, the newspapers followed up with subscription offers at heavily discounted (Black Friday) prices. According to administrative sales data collected by the following February, only three participants had bought a subscription, and only one of them was in the treatment group.

Taken together, our study presents good news, bad news, and a policy conundrum: while professional newspapers promote information and positive political attitudes, they are scarcely demanded even when consumers reacted positively to exposure. The results suggest that policies that increase access to high-quality news could produce positive political spillovers, but which policies can promote consumption remains uncertain. The results imply temporary subsidies need not generate sustained private demand.⁶ Revenue-sharing between social media and newspapers could allow outlets to relax paywalls, but such arrangements have not yet progressed in the market. Public investment in high-quality, non-partisan media could substitute for other market solutions, but political capture could make such a policy backfire (Besley and Prat, 2006).

This study has two important limitations that require a careful interpretation of its results. First, it is unknown whether results of access through subscriptions would extend to a representative sample of citizens, as willingness to engage with the news plays a key moderating role. Second, the three-month subscription period represents a short-run effect. Longer exposure to newspapers can have different effects, potentially varying with newspapers' slant and ideological congruence with consumers.

We contribute to the literature studying the effects of media consumption and subsequent demand for media. Using both observational and experimental data, this literature has studied the effects of newspapers (Gerber et al., 2009; Gentzkow et al., 2011; Cagé, 2020; Hopkins and Gorton, 2024; Briole et al., 2026), television in the United States (Gentzkow, 2006; DellaVigna and Kaplan, 2007; Martin and Yurukoglu, 2017; Broockman and Kalla, 2025; Bursztyn et al., 2023), foreign outlets in China (Chen and Yang, 2019), and social media in the United States, Russia, and Brazil (Enikolopov et al., 2020; Allcott et al., 2020; Levy, 2021; Nyhan et al., 2023; Guess et al., 2023a; Ventura et al., 2025).

First, our study complements the literature on the effects of newspapers on adults by achieving high subscription take-up and identifying the effects of direct exposure among consumers willing to engage: the marginal consumers of news subscriptions. Studies with similar designs have found limited take-ups of local newspapers in the United States (Gerber et al., 2009; Hopkins and Gorton, 2024; Trexler, 2026). Second, we report a substitution away from social-media and toward direct professional-news consumption, and subsequent effects in information, trust, and political attitudes, without a concurrent increase in affective polarization. The pairing of treatment effects with decreases in social media consumption complement recent studies of social media that find limited effects of reducing exposure, including during the Brazilian 2022 elections (Nyhan et al., 2023; Ventura et al., 2025). Second, in contrast to other free-subscription experiments, our treatment allows participants to choose among outlets instead of randomly allocating outlets, providing a naturalistic subsidy to direct consumption. Third, we can measure post-subsidy take-up with administrative data to answer whether demand for news can be sustained given an initial shock, and find null increases in demand. This contrasts with interventions among Chinese students or French high-school students, which

⁶A related implication of our finding is that marketing strategies that lower initial prices of newspaper subscriptions are likely to generate revenue through forgotten subscriptions rather than experimentation (Einav et al., 2025).

found increases in subscription take-up after the intervention (Chen and Yang, 2019; Briole et al., 2026).

The remainder of the article follows the structure outlined below. Section 2 discusses the conceptual motivation for the study. Section 3 details the experimental design, sample, and data. Section 4 describes the treatment implementation and first-stage effects on news consumption. Section 5 presents the treatment effects on political outcomes. Section 6 discusses the heterogeneous treatment effects. Finally, Section 7 concludes.

2 Direct and indirect consumption of professional news

Much of the political information citizens receive reaches them indirectly through social networks and personal connections, not from newspapers themselves, whose content is often behind paywalls. Yet professional reporting supplies a large share of what circulates through those channels. The empirical question is what transmission does to that information. If it distorts content, indirect exposure can lower information quality and weaken trust in the organizations that produced the reporting; direct access may then improve factual knowledge and political attitudes.

Direct access may instead be largely redundant. In many settings, a small number of individuals acquire useful information and diffuse it to others as “influencers” in a two-step flow of communication (Lazarsfeld et al., 1948; Galeotti and Goyal, 2010). Social-media diffusion, however, is not neutral. Algorithms (Guess et al., 2023a), selective filtering (Guess et al., 2023b), political motivations, and social influence (Moehring and Molina, 2026) shape exposure and can degrade information quality (Vosoughi et al., 2018). Political actors may also have incentives to attack professional newspapers, and consumers without direct access are less able to assess whether such criticism is warranted. Indirect transmission can thus substitute imperfectly for direct readership while also reducing trust in professional outlets.

These frictions create a role for direct newspaper consumption in both information and political attitudes, including democratic values. Professional outlets typically produce higher-quality reporting than free content from semi-professional or independent media, even after accounting for slant.⁷ They may also have incentives to defend democratic institutions that protect their profits (Gehlbach and Sonin, 2014; Szeidl and Szucs, 2021). Whether either advantage survives indirect exposure through social networks is ultimately empirical.

⁷Professional newspapers are subject to fair criticism. Even so, they arguably produce higher-quality news because: (i) they staff large newsrooms and maintain on-the-ground bureaus, giving them superior access to primary information; (ii) they impose formal editorial safeguards—bylines, multi-layer fact-checking, rights of reply, and “both-sides” coverage—that reduce the probability of errors; and (iii) they face reputation incentives because their brands, subscription revenues, and prospects for journalism awards are jeopardized by demonstrable inaccuracies. Together, these features make the average report from a paywalled or otherwise professional outlet higher quality than content produced by peer-to-peer networks, semi-professional creators, or partisan aggregators.

3 Experimental design

We study the effects of direct consumption of professional newspapers in the 2022 Brazilian elections. Our experimental design allows us to estimate causal effects of direct consumption of professional news by randomly assigning participants to receive a three-month free subscription to a newspaper. Participants could choose their preferred newspaper among a selection of five of Brazil’s largest professional newspapers, covering right, left, and centrist slants. This choice is meant to approximate market choices and provide naturalistic variation. We measure news consumption patterns and other outcomes through a combination of surveys and administrative browsing datasets provided by the newspapers. Figure 2 displays the experimental design along with sample sizes and the choices of news outlets by the participants.

3.1 Treatment groups, assignment, and implementation

We randomly assigned participants to three treatment groups: control (*C*), free newspaper subscription (*S*), and subscription plus monetary reading incentives (*SI*). Our assignment creates different levels of subsidy intensity to measure causal effects of access to professional newspapers. As participants were able to pick a newspaper from a pre-selected list, a natural interpretation of our treatment is a targeted subsidy: if prices for these news outlets decreased, more citizens would be able to purchase these outlets, and they would do so according to their own preferences.

Participants were assigned to each group with probabilities 55%, 5%, and 40%, respectively. The random assignment was stratified to ensure additional balance. Using baseline data, we stratified users by their choice of outlet, political position, consumption of news, consumption of social media, and news knowledge. These five characteristics were also pre-registered for heterogeneous treatment effects, ensuring balance in the subsets of the data.

Treatment implementation. At the end of the baseline survey, we provided participants in the *S* and *SI* groups a coupon code for a free three-month digital subscription. Coupons were delivered through both email and WhatsApp, along with instructions on how to use the coupons provided by the newspapers. Once participants created their accounts, they were considered regular users by the outlets: they received promotional emails, newsletters, and calls to download apps and follow social media pages as part of the outlets’ typical consumer retention strategies. To further promote usage, we also sent instructions directing participants to download the app and sign up for newsletters.

Reading incentives and comparison between *S* and *SI*. During the first week of the study, we surprised participants in the *SI* group with notice of a three-question quiz every two weeks on content specific to their chosen outlet.⁸ We informed them that good performance in the quizzes would be rewarded with additional monetary incentives. In expectation, the average participant would obtain a payment of US\$4 — in comparison, the value of the three-month subscription was between US\$8 and US\$12. Thus, the encouragement is best regarded

⁸The news quizzes were short, relatively easy, and tailored to each newspaper, comprising questions about the newspapers’ content covered in the past two weeks.

as a nudge rather than a substantial monetary incentive, similar to Chen and Yang (2019).

The extra incentives were used primarily to guarantee a relevant first stage, motivated by ample evidence of low demand for news subscriptions (e.g., Hopkins and Gorton, 2024; Batista Pereira et al., 2023). Therefore, we allocated a much larger fraction of participants to the subscription plus incentives group *SI* (about 40% compared to 5%).

3.2 Sample recruitment and description

To explore the contrast between social media and professional news consumption, we recruited our sample through ads on social media platforms. The eligibility criteria were: (i) being a Brazilian citizen; (ii) being 18 years or older; (iii) having WhatsApp installed on their phone⁹; (iv) not currently subscribing to any news outlets; (v) declaring an interest in using the free three-month subscription. The recruitment ad and the social media profile associated with it both mentioned an “academic news study”, so subjects clicking the link already knew the study was related to news.¹⁰

The baseline analysis frame comprises the 1,658 respondents who completed the baseline survey. We identified 23 respondents sharing an identical IP address and exact GPS coordinates with another respondent, and retained them in the sample.

The recruitment strategy is motivated by evidence of low demand for subscriptions and we therefore aimed to recruit a sample of “marginal consumers” of digital news subscriptions: participants who expressed interest in having a news subscription, but had decided not to pay for a subscription. This “marginal news consumer” is the immediate sub-population of interest for policies that encourage news consumption or subsidize spending on news outlets. The trade-off for the eligibility criteria, along with online recruitment, is a lack of representativeness. Table A.1, columns 1 and 2, displays the summary statistics for our sample. The sample skews toward certain demographic groups: women, older individuals, left-wingers, and more educated individuals.¹¹

As participants declared interest in the news subscription, we may expect them to be more news-savvy than the overall population. Participants report a strong preference for consuming news on social media. The average self-reported news consumption is about 62 minutes per day, with social media usage declared to be about 115 minutes per day. However, self-reported measures of news consumption are known to be biased in levels: according to the most recent evidence from Pew Research Center (2020), some self-reported measures may be biased upward by a factor of two, while prior research suggested an upward bias by a factor of three (Prior, 2009).

⁹According to Datafolha, the main Brazilian pollster, WhatsApp was used by 92% of those interviewed in March 2022 and is by far the most used social app among Brazilians. The numbers vary by years of schooling: among those who have completed only elementary school, 82% use WhatsApp; among those with high school education, it is 95% and for higher education it is 98%. Source: <https://www1.folha.uol.com.br/tec/2022/07/94-tem-conta-em-alguma-rede-social-whatsapp-lidera-com-92.shtml>

¹⁰See Appendix Figure A.1 for an example of an ad.

¹¹See Figure 2 for more details on the experimental design.

3.3 Data and timeline

We measure treatment implementation and effects through combined newspaper browsing data, baseline and endline surveys for political attitudes, and periodic WhatsApp surveys tracking news consumption and knowledge.

Main surveys. The main survey dataset combines information collected in the baseline and endline surveys. Measurements of the pre-registered outcome variables are collected in both surveys.

Phone surveys. The phone surveys dataset contains data collected through WhatsApp surveys.¹² It contains five surveys about news consumption behavior, which were administered every two weeks, and two news knowledge quizzes designed around key political events: the September 7 Independence Day parades and the second round of the general elections. We label the samples used for each measurement as the “endline survey sample”, and the phone surveys as “September 7 survey sample” and “pre-elections survey sample”.

Administrative newspaper data. The newspapers participating in this study provided administrative data on coupon redemption, online news consumption, and subscription purchases, allowing us to directly measure treatment take-up, consumption patterns over the course of the study, and post-intervention demand for subscriptions.

Timeline. Figure 1 summarizes the study timeline. Recruitment for the study took place between August 21 and August 29. The recruitment period coincided with the start of the official campaign period for the 2022 elections, when advertisements on television and social media began. We conducted the first news knowledge survey between September 9 and September 11, a few days after the September 7 Independence Day parades, a key political event. The first round of the elections took place on October 2, when Congress members were elected and second-round contests were determined. In particular, the presidential race proceeded to a second round on October 30. We conducted a phone survey shortly before, between October 28 and 29, asking questions about news covering the electoral period. Finally, participants could answer the endline survey between November 26 and December 3. Figure 1 below displays the timeline of the study along with the main surveys conducted.

3.4 Outcome variables

Following the pre-analysis plan, we divided the main outcomes into four families: (i) news consumption and opinions; (ii) information knowledge; (iii) political attitudes; and (iv) democratic preferences. Indices for each family of survey outcomes are built using variance-covariance weighting (Anderson, 2008). Variable measurements are described in detail in Appendix C. Deviations from the pre-analysis plan are noted in Appendix F.

The first family measures outcomes directly related to the newspapers and the press. Regarding consumption, we obtain measures of daily news consumption and news diet across

¹²Surveys were administered directly through WhatsApp. Participants received a message with a button to begin the survey. Those who began their responses and failed to respond within one minute after being sent the question had their surveys terminated.




Figure 1: Study timeline, 2022

Notes: Horizontal position is proportional to calendar time; ticks mark the first day of each month. Shaded bars mark the fielding window of each survey wave; dots mark political events; the two rounds are those of the 2022 general election, on the dates set by the Superior Electoral Court (TSE).

media (news websites/apps, TV, social media, etc.), and whether the participant decides to purchase a subscription after their free subscription ends. We also ask about trust in their newspaper of choice and opinions about the press and journalists.

The second family measures news and political knowledge through four quizzes. Three were news quizzes, composed of statements about the news cycle of the previous two weeks. The statements were divided into three categories: “true news”, which were true statements; “fake news”, which were selected with the help of fact-checking websites; and “incorrect news”, which were false statements constructed by slightly altering true news statements.¹³ The fourth information quiz focused on political knowledge and was designed to have questions pertaining to the electoral rules and composition of Congress, tying the political knowledge questions to the 2022 general elections. Each quiz is graded by assigning one point to a correct answer, half a point to “don’t know”, and zero points to an incorrect answer.

The third family of outcomes measures political engagement and affective polarization. Political engagement was measured using an index of eight political activities. Seven items were measured on a five-point frequency scale (from “Never” to “Very often”); the campaign-donation item was measured on a four-point scale. Affective polarization was primarily measured with respect to Lula and Bolsonaro, using the “feeling thermometer” that captures whether participants feel positively or negatively with respect to the presidential candidates.

The fourth family measures outcomes related to democracy.¹⁴ Regarding elections, as a primary outcome we measure participants’ trust in the electoral system, and as a secondary outcome, we measure support for a proposal from the Bolsonaro government that suggested changes in the electoral system. To capture pro-democratic preferences, we adapt the support-for-democracy measure from Graham and Svolik (2020). Participants are shown six statements describing actions by a hypothetical politician that could be interpreted as undemocratic, such

¹³To choose the most widely circulated fake news, we had assistance from *Estadão Verifica*, which is the in-house fact-checking agency of one of the news outlets we partnered with, *Estadão*. Despite being specific to a newspaper, *Estadão Verifica* is part of a broader alliance of fact-checking agencies that are reported by many news outlets.

¹⁴To reduce normative judgments, we do not assume that greater trust in the electoral system is necessarily a pro-democratic attitude.

as banning opposition protests or supporting military interventions. We then ask if the participant would be less or more willing to vote for that politician. In this case, a lower willingness-to-vote is interpreted as a preference to punish candidates who deviate from democratic norms.

3.5 Internal validity: balance, attrition, and experimenter demand effects

Balance. Table A.1 displays the main balance tests. Columns 3, 4 and 5 present the averages of demographic and media consumption variables for the Control (*C*), Subscription (*S*), and Subscription+Incentives (*SI*) groups. The p -values of the balance tests displayed in columns 6 and 7 are obtained from a regression of each variable on treatment dummies and strata fixed effects. Rejection rates of differences between the main Subscription+Incentives and Control groups in 22 key demographic and media consumption variables are within expectation: one is significant at the 5% level, and one is marginally significant. The differences are also quantitatively small.

Figure A.2 shows balance across four samples (baseline, endline, and the two phone surveys). In this figure, we compare the *SI* group with the control group. All four samples are similarly balanced, with p -values of joint significance F -tests between 0.29 and 0.63. The magnitudes of differences are also small: no absolute average difference in covariates is larger than 0.11 standard deviations.

Attrition. The endline survey sample is composed of 1,363 participants, representing a completion rate of 82%, which is within the usual range of completion rates in field experiments (Ghanem et al., 2026).¹⁵ Completion rates for the pre-election phone survey and the September 7 phone survey were 73% and 80%, respectively.

We find evidence of differential attrition. Table A.15 compares Control (*C*), Subscription (*S*), and Subscription + Incentives (*SI*): the *S* and *SI* groups are about 12 to 15 pp. less likely to respond to post-baseline surveys ($p < 0.01$, columns 1, 3 and 5). Differential attrition in the *S* and *SI* groups is driven almost entirely by the roughly 30% of treated participants who did not redeem their coupons (columns 2, 4, and 6).¹⁶ The difference is already present in the first phone survey, conducted roughly one week after baseline, rather than compounding over time (columns 1 and 2).

Figure A.2 presents an ex-post balance test on baseline covariates and shows that the post-baseline samples remain balanced. In Appendix D we examine attrition bias following Ghanem et al. (2026) and report bounds following Semenova (2025). Baseline outcomes and demographics predict neither attrition nor differential attrition. However, treatment take-up can correlate with the demand to consume news. We compare compliers and non-compliers on baseline measures of *expected* extra consumption induced by the treatment. Compliers and non-compliers do not differ in their expected consumption change. Given this evidence, we conjecture that

¹⁵Similar studies employing online recruitment are rarer in Brazil, but a recent study with a similar setup reported 65% and 73% completion rates for two different waves when using a survey company instead of recruiting through social media (Batista Pereira et al., 2023).

¹⁶Subscription take-up is discussed specifically in Section 4.2.

sign-up and coupon-activation frictions caused some participants who failed to activate their free subscription to leave the study.

Experimenter demand effects. Because most outcomes are self-reported, treatment effects could partly reflect what participants believed we wanted to hear, even though the survey is fully online and we never directly interacted with participants. Two pieces of evidence limit this concern. First, at endline 69% of participants believed the study promoted no particular view of news outlets and 71% believed it promoted no political side. The *SI* group was 8.2 pp. more likely than the control group to believe we were promoting a positive view of outlets ($p < 0.01$; Appendix Table A.4), consistent with the reading incentives, while the *S* group did not differ (-1.3 pp., $p = 0.85$). Beliefs about political sides are balanced across arms. The numbers suggest a limited scope for experimenter demand effects to influence participants' responses.

Second, the outcomes most exposed to experimenter demand effects are those directly related to news consumption, the object of the study. We find no evidence that participants inflated their self-reported coupon redemption: 56 redeemed a coupon but reported they did not, while only 24 participants claimed a coupon we could not match in administrative data. Relatedly, re-estimating the news-diet effects on the subsample who did not believe we were promoting a view of the news leaves the estimates, including the increase in outlet website and app consumption, unchanged (Figure A.3).

3.6 Selection of the professional news outlets and their coverage

The selection of professional news outlets attempted to cover the whole ideological spectrum in order to provide enough choices to meet participants' preferences. Outlets had to have a paywall on their content in order for the free subscription to be meaningful. We partnered with three of the four leading outlets in the country, which are large and traditional newspapers with a relatively centrist slant.¹⁷ We also partnered with flagship professional publications on the left and on the right, which are smaller in size but still arguably the largest within their ideological slants. Although the ideological position of newspapers can be debated, we based our description on their support for or opposition to the presidential candidates. Table 1 presents information about the chosen newspapers, including circulation numbers and web traffic data.

Based on their editorials, there was some convergence in topics discussed leading up to the election: polarization, political violence, civil unrest, mistrust in elections, and risks to democracy. The larger publications tended to criticize both candidates, but overall were more vocal in their opposition to Bolsonaro. Given the electoral context and our focus on democratic attitudes, the editorials of the right-wing publication are of special interest. On the one hand, *Gazeta do Povo* was critical of judicial decisions over the course of the electoral process, in particular those restricting certain speech, regarding the decisions as censorship. On the other hand, the newspaper denounced attempts to discredit the electoral system in a July editorial. After the elections concluded, the publication was critical of the protesters and of the civil

¹⁷The other newspaper in the top four, *Folha de São Paulo*, was unable to participate in the study due to technical constraints on coupon distribution.

unrest that took place, being explicitly critical of Bolsonaro in both cases. Overall, the newspapers’ professional reporting and editorials did not seem to promote misinformation about the electoral system or anti-democratic attitudes.

Table 1: Summary of selected news outlets

Outlet	Year founded	Circulation (2020)	Political orientation	Similarweb traffic data	
				Visits (3 months)	Ranking (News and Media)
Carta Capital	1994	26,000	Left	3.3M	61
Estadão	1875	213,000	Center	30.5M	14
Gazeta do Povo	1919	80,000	Right	9.8M	27
O Globo	1925	370,000	Center	55.9M	7
Valor Econômico	2000	131,000	Center	8.3M	31

Note: Circulation numbers are from the non-profit organization Instituto Verificador de Comunicação (IVC), whose objective is to provide reliable data on media circulation in Brazil. One exception is Gazeta do Povo, which is not associated with IVC, and thus we rely on its self-reported number of subscriptions. We also use data from Similarweb’s News & Media category for online traffic information; the numbers were checked in March 2023.

3.7 Empirical specification for treatment effects

This section specifies our main empirical specifications to estimate treatment effects. Let Y_i denote an outcome, and let Subscription_i and $\text{SubscriptionIncentives}_i$ be indicators for assignment to the S and SI groups, respectively.

The first estimating equation is

$$Y_i = \alpha + \beta_S \cdot \text{Subscription}_i + \beta_{SI} \cdot \text{SubscriptionIncentives}_i + X_i' \delta + \varepsilon_i, \quad (1)$$

where the control vector X_i is chosen separately for each outcome using the post-double-selection method of Belloni et al. (2014) over a set of baseline variables. In the control matrix, qualitative variables and variables with five or fewer unique values are recoded as indicator variables for each value. When variables lack a natural scale, we standardize Y_i by subtracting the Control group mean and dividing by its standard deviation. Indices are divided by the full-sample standard deviation. For Equation (1), standard errors are heteroskedasticity-robust (Abadie et al., 2023). X_i always includes dummies for the chosen news outlets.

Our main coefficient of interest is $\hat{\beta}_{SI}$, as we designed this group to be the largest treatment group and the one with the strongest first stage. All figures presenting estimates in Section 5 and Section 4 display the estimates $\hat{\beta}_{SI}$. We report estimates $\hat{\beta}_S$ in companion tables in the Appendix.

To estimate heterogeneous treatment effects, we augment Equation (1) in two ways. First, let H denote the variable for which we estimate heterogeneous effects. If H is binary, we estimate the following equation based on differences:

$$\begin{aligned}
Y_i = & \alpha + \beta_1 \cdot \text{SubscriptionIncentives}_i + \beta_2 \cdot \text{SubscriptionIncentives}_i \times H_i \\
& + \beta_S \cdot \text{Subscription}_i + \beta_{S,H} \cdot \text{Subscription}_i \times H_i \\
& + \gamma H_i + X_i' \delta_1 + H_i \cdot X_i' \delta_2 + \varepsilon_i,
\end{aligned} \tag{2}$$

where X_i is selected by post-double selection without forcing addition of the outlet dummies. β_2 represents the average difference in treatment effects between participants with $H = 1$ and $H = 0$. The Subscription terms keep the S arm in the sample with its own main and differential effects, which we do not report.

If H takes three or more values, instead of an equation based on differences, we estimate a non-saturated equation:

$$\begin{aligned}
Y_i = & \sum_{h \in H} \beta_h \cdot \text{SubscriptionIncentives}_i \times \mathbb{1}[H_i = h] + \sum_{h \in H} \beta_{S,h} \cdot \text{Subscription}_i \times \mathbb{1}[H_i = h] \\
& + \sum_{h \in H} \gamma_h \cdot \mathbb{1}[H_i = h] + X_i' \delta + \sum_{h \in H} \mathbb{1}[H_i = h] \cdot X_i' \delta_h + \varepsilon_i,
\end{aligned} \tag{3}$$

4 Treatment implementation and first stage

In this section, we first examine the choice of newspapers by participants, treatment take-up, and consumption patterns among the treatment group using administrative data provided by the newspapers. We then estimate the intervention's first stage by comparing measures of news diet across the control and treatment groups.

4.1 Choice of news outlet

How do participants choose outlets, and how does their choice affect the interpretation of results? Prior literature identifies perceived ideological alignment and accuracy concerns as two strong drivers of newspaper demand (Gentzkow and Shapiro, 2010; Chopra et al., 2022). If participants choose in line with these standard preferences rather than, say, using the free subscription to sample counterattitudinal outlets, we can interpret the treatment as a subsidy that delivers naturalistic price variation with a conventional demand response. The data support this interpretation.

During the baseline survey, we asked participants how much they trusted each outlet and their perceived ideology of each outlet (according to a seven-point scale from extreme left to extreme right).¹⁸ We assume accuracy concerns, or quality, are proxied by trust, and that ideological distance, measured as the difference between a participant's own reported ideology

¹⁸We also allow an "unfamiliar" option if participants are unable to provide an opinion about trust, ideology, or both. Unsurprisingly, the smaller outlets Gazeta do Povo and Carta Capital, which position themselves as right- and left-wing, respectively, are the most unfamiliar to participants, along with Valor Econômico, the smallest, economy-focused outlet among the three largest.

and their perceived ideology of each outlet, is a proxy for ideological alignment.

Participants chose like-minded and trusted outlets. Among those who reported both their own ideology and their chosen outlet’s, 28% of choices are in complete ideological alignment and a further 50% are within two units. On the four-point trust scale, 86% chose an outlet they trust “a great deal” or “a bit”, and only 4% chose one they were unfamiliar with. Because trust and ideological distance are strongly correlated, we estimate a multinomial logit of choice on both (Appendix Table A.2): conditional on trust, each unit of ideological distance lowers the relative probability of choosing an outlet by about 20–25%. Among equally trusted outlets, the more aligned one is more likely to be chosen.

4.2 Treatment take-up and consumption of news from chosen news outlets

We use administrative data to measure participants’ usage of coupons and their consumption of news from their chosen outlets.

Treatment take-up. 526 out of 751 coupons were redeemed across both treatment groups, amounting to a 70.0% take-up. This is a relatively high redemption rate, in line with our decision to screen out participants who were uninterested in the subscription.¹⁹ Coupon redemption ranged from 63% to 74% across outlets.

Beyond redemption, treated participants also took active steps toward consumption. In the endline survey, among those who reported using the coupon, 59% reported downloading the app, 55% signed up for newsletters, and 44% reported following the outlet on social media. Concerning the consumption medium, cellphone apps were the preferred method of 37% of participants, followed by computer browsers (32%) and mobile browsers (30%).

Outlet consumption data. Each newspaper’s browsing data record a URL and a timestamp associated with the account’s coupon. Internal collection methods vary. Browser data typically use Google Analytics, while each newspaper’s apps capture these data in different ways. A key advantage of these data is coverage: it captures both apps and browsers, on phones and computers, with no opt-in from tracking extensions typically used to measure news consumption.²⁰

The data have two important limitations. First, Carta Capital had a soft paywall where login is not always required to read an article. Similarly, O Globo shares a login system with a broader entertainment platform. As a result, the match between coupon codes and browsing data is incomplete for these two outlets: matching rates of Carta Capital and O Globo are 39% and 35%, respectively, compared to about 86% for the others. Second, time-on-page is inferred from timestamp differences, so inactivity, single-URL sessions, and in-app browsing introduce

¹⁹As a comparison, other studies examining representative samples found much lower redemption rates. Using a representative sample of Brazilian citizens and offering coupons for free subscriptions of one national-level outlet, Batista Pereira et al. (2023) had an 11.4% redemption rate. Hopkins and Gorton (2024) offered Pennsylvania residents free subscriptions to leading state newspapers, but found only a 1.7% redemption rate.

²⁰Tracking extensions on Google Chrome are a common alternative, but they require an opt-in decision and cover only computer browsing, not smartphones or apps, which is limiting given participants’ preference for mobile consumption.

measurement error in levels. In addition, in a browsing test we found that on average the data collection misses 6.1% of unique URLs browsed. In Appendix Section C.1, we discuss harmonization of the dataset across sources, specific issues, and how we define the browsing variables.

Our three comparisons are robust to the data limitations under mild assumptions. If missingness is independent of date, we can compare consumption over time. If it is independent across individuals, we can examine the distribution of consumption among participants. And if it is independent of treatment assignment, we can test whether reading incentives worked by comparing the *S*-versus-*SI* difference. The browsing test supports these assumptions.

Consumption of news over the study period. Figure 3 presents daily news consumption (in minutes per day) over the twelve weeks of the study. The gray shaded areas refer to three key events during the period: the September 7 parade and the two rounds of the elections. Panel A displays average daily news consumption, averaged both over all treated participants and over those present in the browsing dataset. Panel B presents average daily news consumption conditional on consuming some news that day, along with a second axis displaying the number of participants consuming some news that day.

We observe consumption peaks during the September 7 parade and the two electoral rounds. Furthermore, after the election, consumption steadily drops as the news cycle turns to mostly cover the transition of government, both in total news consumption time and in the number of people browsing some news on the platforms.

Distribution of news consumption. Table 2 presents summary statistics for the browsing data at the participant level. The key takeaway is that consumption varies widely: between the 25th and 75th percentiles, there is a wide variation in the number of days browsing (2 to 23), total URLs consumed (3 to 37), and browsing duration (347 to 7,690 seconds). This pattern is consistent with news consumption being highly skewed, with a minority of participants accounting for most consumption.

Differences between treatment groups. The extra incentives made participants more likely to initiate consumption of the outlet, while consumption conditional on browsing was similar across the two groups (Table 3).²¹ Coupon redemption rates do not differ (column 1), but the *SI* group is 11.7 pp. or 35% more likely to have some browsing activity (column 2, $p = 0.02$). Conditional on having some browsing activity, we find no differences in the number of days browsing (-0.6 days, $p = 0.88$) or in total URLs browsed (3.5 URLs, $p = 0.70$; columns 3 and 4). Setting consumption to zero for participants without any browsing data, the differences are positive but imprecisely estimated (2.2 days, $p = 0.16$; 5.1 URLs, $p = 0.17$; columns 5 and 6).

²¹We analyze differences between the *S* and *SI* groups in both extensive and intensive margins of consumption using the administrative browsing dataset: (i) coupon redemption rates; (ii) presence in the browsing dataset; (iii) overall browsing, conditional on being present in the browsing dataset; and (iv) overall browsing. Measures (i) and (ii) are extensive margins, while measure (iii) is the intensive margin. Measure (iv) adds both margins by coding browsing metrics for those without browsing data to zero. For each measure we regress Y_i on an SI_i indicator and controls X_i , so differences are captured using SI_i as the baseline group; as in our other regressions, we select controls through post-double-selection and report heteroskedasticity-robust standard errors.

5 The effects of increasing access to professional outlets

Shifting news consumption from secondary sources toward professional newspapers, without increasing total news consumption, increased news knowledge and trust in the press, strengthened support for democratic norms, and increased political engagement without increasing affective polarization. However, it generated virtually no lasting demand for paid subscriptions.

Our figures display the treatment effects for the *SI* group. Results for both groups are displayed in Tables A.5–A.12. Pooling the two treatment arms into a single indicator leaves the family indices essentially unchanged (Appendix Table A.14).

5.1 Changes in news diet

Providing free subscriptions and incentives to consumption changed news diets by pushing consumers toward news from websites/apps of outlets, without significantly changing total news consumption. This increase in consumption induced a substitution away from consuming news on social media, with a suggestive complementarity with television news. We therefore interpret the treatment as a shift across sources rather than an increase in total consumption.

Total news consumption barely changes at endline. Figure 5 reports an increase of about 0.3 minutes (0.4% of the control-group mean of 70.37 minutes), or 0.01 standard deviations, and we cannot reject the null ($p = 0.88$). The 95% confidence interval ranges from a decrease of about 3.7 minutes to an increase of about 4.3 minutes. Participants nevertheless report following politics more closely: the increase is 0.10 standard deviations ($p = 0.04$).

The experiment shifted the composition of news consumption away from social media and toward website/apps of newspapers. We estimate a strong and precise increase in the consumption of news from news outlets’ websites/apps (0.20 s.d., $p < 0.01$), while social media consumption intensity fell by 0.08 standard deviations ($p = 0.09$). We see a suggestive spillover into television news consumption: cable news, local TV, and national TV all increased, though only local TV is marginally significant (0.07 s.d., $p = 0.08$). News consumption from video platforms and radio/podcasts was largely unchanged. In a measurement of the main medium of consumption, the treatment decreased “social media” as the main medium by 4.5 pp. (15.5% of the control group mean, $p = 0.05$), and increased website/apps and TV, though imprecisely.

The temporal trends measured over the course of the study broadly agree with endline results and actually show that total consumption in the *SI* group was sometimes lower than in the Control group. Figure 4 reports the per-wave *SI*–Control differences for the same outcomes, each estimated as in Equation (1): the website/app gap is positive across the post-baseline waves, the social-media gap is negative, and total consumption shows no systematic increase.

5.2 News and political knowledge

The treatment raised average news knowledge, especially during the electoral period. We measure news knowledge through three quizzes on the political news cycle, two fielded during

the campaign and one at endline. Combining the three quizzes into a variance-covariance weighted index, we find a treatment effect of 0.10 standard deviations ($p = 0.06$). General political knowledge does not move in the same direction: the estimate is negative and not statistically distinguishable from zero, at -0.01 standard deviations ($p = 0.85$).

Each news knowledge quiz covers the news cycle of the previous two weeks. Participants classify each statement as true or false, with an “I don’t know” option. The statements are of three types: “true news”, “incorrect news” (real headlines slightly altered into factually incorrect statements), and “fake news”.²² We measured political knowledge through an endline-survey quiz on general political facts, such as the number of Senate and Chamber of Deputies seats by state, which were the largest parties in Congress, and who were the current presidents of the Senate and the Chamber of Deputies.

Figure 6 presents treatment effects of increasing access to professional news outlets on news knowledge based on Equation (1).

The first news quiz contained six questions and was conducted two days after September 7, Brazil’s Independence Day. It focused on news about the parade held in Brasília by the government, a state-sponsored event considered a political rally for candidate Jair Bolsonaro. The study began roughly one week before the parade. The quiz used only “true news” and “incorrect news”, as few explicitly fake stories circulated around the event. The treatment group had a quantitatively relevant but noisy increase: we find an effect of roughly $+0.07$ standard deviations on the quiz score ($p = 0.20$).

The second news quiz contained nine statements and was distributed two days before the second round of the elections. The statements therefore focused on the news cycle of the last week of the presidential campaign, when news consumption was at its peak. Participants in the treatment group performed better than the control group, with a difference of $+0.13$ standard deviations, significant at the 5% level ($p = 0.03$).

The endline quiz was fielded about one month after the presidential election. By then the news cycle centered on executive transition teams and the institutional response to protests and road closures that began after the elections.²³ The overall score is not statistically distinguishable between treatment and control: $+0.04$ standard deviations ($p = 0.37$). The knowledge quiz fielded at endline covers general political facts and current affairs, with most questions tied to the election context. The estimate is negative and not statistically distinguishable from zero, at -0.01 standard deviations ($p = 0.85$).

The differences in results across surveys are consistent with consumption patterns. Figure 3 shows that news consumption collapsed after the elections, which can explain the weaker effects in the post-electoral quiz. Consumption instead spiked during the September 7 and election periods, where we find positive treatment effects.

Treated participants also reduced their frequency of “don’t know” responses to incorrect ones. If we instead attribute 0 rather than 0.5 points to “I don’t know” responses, the treatment effect

²²Fake news are stories that actually circulated during the campaign, selected from the fact-checking websites *Projeto Comprova* and *Estadão Verifica*, with the assistance of journalists from *Estadão Verifica*.

²³By the end of November, most protests had ceased.

risers to 0.3 standard deviations ($p < 0.01$).

5.3 Opinions about newspapers and the press

Increased exposure to the outlets strongly increased trust in those outlets, and also led to spillovers in trust in journalists and the media as a whole. The effects are much larger for those with low baseline trust or who were unfamiliar with the outlet, and the increase in trust was not led by a change in perceived ideologies, which were unmoved by the treatment. The sign of this effect was not obvious *ex ante*: readers could have updated negatively after seeing slant or selective emphasis, or not updated at all if beliefs about well-known outlets were already fixed.

The treatment has a large and precise effect on trust in the chosen outlet (Figure 7): a 0.30 standard deviation increase ($p < 0.01$). Splitting the results by baseline trust, the effects are significant both for participants who reported medium or high trust at baseline (0.23 s.d., $p < 0.01$) and for those who reported low trust or reported not knowing the outlet they chose (0.71 s.d., $p < 0.01$).²⁴ Outlets' perceived ideology was mostly unchanged; only those who did not know the outlet's ideology at baseline changed their perceived ideology at a higher rate than the control group (0.28 s.d., $p = 0.05$).

Trust in other news sources moves in the opposite direction, though the estimates are less precise. Trust in news from friends and family falls by 0.08 s.d. ($p = 0.10$) and trust in social media news by 0.05 s.d. ($p = 0.38$). The contrast with the increase in trust in the chosen outlet and the press suggests participants updated on the quality of newspaper news relative to other sources, rather than becoming more trusting of news in general.

An index of opinions about journalism and the press increases by 0.09 standard deviations ($p = 0.02$). All three components move in the same direction: a noisy 0.04 standard deviation increase in confidence in the press ($p = 0.36$), and 0.09 ($p = 0.05$) and 0.10 ($p = 0.02$) standard deviation increases in agreement that journalists provided fair coverage of elections and of politics. This is consistent with spillovers from consumption of the outlet into the broader press, as professional newspapers are important representatives of the media as a whole.

Our results provide one interpretation of the declining confidence in the press documented in many countries over the last decade, which has coincided with a decline in the consumption of newspapers (Newman et al., 2022). The natural direction of causation is that lower confidence reduces demand for professional outlets. Our results establish the converse, so the shift toward digital consumption may itself have reduced exposure to flagship outlets and diminished confidence in the press as a whole. It is also aligned with an increase in criticism of journalists by key politicians across the globe. In our context, 65% of Brazilian respondents report hearing some criticism of journalists or the media “quite often” or “very often”.

²⁴The results are very similar for those who reported low trust and for those who reported not knowing the outlet.

5.4 Political engagement

The treated group increased its average political engagement relative to control, although the combined evidence is imprecise. Participants in the treatment group report following politics more closely (0.10 s.d., $p = 0.04$). The pre-registered index of eight political activities moves in the same direction but is imprecisely estimated (0.06 s.d., $p = 0.11$).

Figure 8 reports the eight activities and the pre-registered index. Election-specific activities all move positively: watching debates for governor (0.08 s.d., $p = 0.09$) and persuading others to vote (0.08 s.d., $p = 0.09$) are marginally significant; reading candidate interviews yields a similar point estimate (0.07 s.d.) but is imprecisely estimated ($p = 0.14$). The increase in debate watching is in line with a suggestive increase in consumption of news through local television (Figure 5). The other individual items are statistically indistinguishable from zero.

Turnout in the elections was pre-registered as a secondary outcome because of compulsory voting. Indeed, 91% of participants in the control group reported voting. Still, effects on turnout are positive in both rounds but imprecisely estimated (0.03 s.d., $p = 0.54$; 0.07 s.d., $p = 0.18$).²⁵ If instead we define turnout as voting in either of the two rounds, the effect is 0.08 standard deviations ($p = 0.13$). Other pre-registered secondary outcomes are displayed in Appendix Table A.13. Only one outcome is statistically significant: the treatment decreased reported trust in government by 0.09 standard deviations ($p = 0.02$).

5.5 Affective polarization

Affective polarization does not change due to the treatment (Figure 9). Given the large number of political parties with Congress representation in Brazil, our primary measure is the affective polarization between the two main candidates, Bolsonaro and Lula. We also measure affective polarization between the Supreme Court and Bolsonaro, which were increasingly antagonistic during the COVID-19 pandemic, and between the “Left” and “Right” political fields, one measure concerning supporters and another concerning politicians. Results are null across all measures.

The null contrasts with evidence that television and social media can increase polarization (e.g., Martin and Yurukoglu, 2017; Levy, 2021; Broockman and Kalla, 2025; Bursztyn et al., 2023), and is consistent with professional newspapers being a different medium: written, largely “hard” news reporting, with opinion pieces in a separate section and limited scope for the casual language and imagery of television and social media (Kalmoe et al., 2018). We can rule out a meaningful net increase in polarization from allowing participants to choose ideologically congruent outlets, even though the treatment also raised trust in those outlets and engagement with politics. Because the design does not isolate moderated language, information, or ideological congruence, we cannot tell whether these channels are individually null or offset one another.

²⁵The second-round turnout estimate corresponds to about 2 percentage points on a base of 91%.

5.6 Pro-democratic attitudes

We ask whether exposure to professional outlets increases pro-democratic attitudes. The professional press has increasingly adopted explicit messaging in favor of democracy and democratic institutions, both in its editorials and in its reporting lines.

First, we measure a hypothetical “willingness-to-vote” for politicians who commit undemocratic actions, employing a simplified version of Graham and Svolik (2020). This measure is motivated by a view that democratic erosion takes place as a slow and incremental process, where politicians make decisions that erode norms. Often, punishing norm deviation through voting is the main avenue to avoid democratic erosion. A low willingness-to-vote captures the disposition of voters to punish politicians who hypothetically commit actions considered to be undemocratic. On average, most of our sample already expressed a low willingness-to-vote: the median response was 0 (“much less willing to vote”, on a 5-point scale) for all but one measure.

The treatment decreases the measure of willingness-to-vote by 0.06 standard deviations ($p = 0.13$), a mild pro-democratic shift. Figure 10 displays the results. Since we framed the questions as undemocratic actions, a lower willingness-to-vote represents a more pro-democratic attitude, so negative estimates are pro-democratic shifts. For the six measures, five are negative, with the exception of disrespecting judicial orders. The journalism-related item, which could plausibly be directly affected by the treatment, is close to zero. As we show in Section 6, results are driven by the non-left-wing participants.

Second, we estimate whether the treatment increased trust in the electoral system. Over the course of the 2022 general elections, the electoral system was a heavily contentious issue. The controversy predates the campaign. In 2021, the government sought a constitutional amendment requiring electronic ballots to print a voting receipt that would be deposited automatically in a physical ballot for auditing, the “auditable printed vote” (*voto impresso auditável*).

There are three pre-registered primary outcomes measuring trust in the electoral system: agreement with a statement that “elections held in Brazil are free and fair”, a 4-point scale asking about overall trust in the electoral system, and a 4-point scale asking about confidence in elections (following the same wording from the World Values Survey). 69.5% of the sample considered the Brazilian elections free and fair at baseline, and 56.7% expressed complete trust in the electoral system. As secondary outcomes, we inspect whether participants are aware of the “auditable printed vote” and their opinions, as this was a hotly debated topic and the main proposal of the Bolsonaro government to change the electoral system.

To provide context, consider a poll by AtlasIntel conducted right before the second round of the elections, which found that 36.6% of the electorate did not trust the electronic voting, with a deep partisan divide: roughly 88% of those who did not vote for Bolsonaro in the first round expressed trust in the ballots, compared to only 9% for Bolsonaro voters.²⁶

As professional newspapers usually publish hard news, they reported both the attacks on the system by the Bolsonaro campaign and the reactions from other political entities. They also reported widely on reaction by the Superior Electoral Court to assure citizens of the integrity

²⁶The poll was registered with the Superior Electoral Court under registry BR-06415/2022. Retrieved from <https://cdn.atlasintel.org/f2bfa8f9-ef18-4b42-9c1f-918b5949f6f5.pdf>.

of elections. In particular, of the newspapers included in this study, most were energetic in their publishing of fact-checking concerning the ballots and in their defense of the electoral system.

The primary trust measures are largely unchanged and precisely estimated. For the variance-covariance-weighted index of the three outcomes, Figure 11 shows a confidence interval that rules out changes of about 0.06 standard deviations. Among respondents who know the proposal, however, support falls by 0.09 standard deviations, significant at the 5% level ($p = 0.04$), while support shifts toward retaining electronic voting (0.09 s.d., $p = 0.02$).

The average effects on trust in elections, however, hide heterogeneity by political position as we show in Section 6: non-left-wing participants reported an increase in trust in the electoral system (0.10 s.d., $p = 0.09$), and this effect differs from the one among left-wing participants at the 5% level ($p = 0.04$).

5.7 Post-experimentation demand for news subscriptions

We find virtually no demand for news subscriptions after the experiment. We measure post-experiment demand through two outcomes: participants' stated interest in a Black Friday subscription offer and their actual subscriptions in the months after the experiment ended. The treatment group was slightly less likely to express interest in the Black Friday offer, although the estimate is imprecise (Figure 12). Administrative sales data are even starker: roughly two months after the experiment ended, only three participants had purchased a subscription, only one of whom was in the treatment group.

One way to interpret the intervention is allowing consumers to experiment with professional news. Experimentation can increase valuation by reducing uncertainty about the product, and repeated consumption may also induce habit formation. If exposure raises subsequent demand, temporary subsidies could therefore generate a positive effect in which increased exposure leads consumers to continue purchasing news after the subsidy ends. Our results provide little evidence for such a mechanism. Thus, promotional rates or welcome offers used to generate entry into subscriptions may generate revenue through mechanisms other than learning the value of the product, such as inattention (Einav et al., 2025).

This is a particularly relevant test because our study sample consists of marginal consumers of news subscriptions: participants were interested enough to join the study and expressed interest in using at least one of the professional outlets we partnered with, but did not already subscribe, primarily because of price (Section 3.2). Our design asks whether sustained exposure can increase demand among a population considerably closer to the margin of purchasing. The answer is negative.

The low post-study demand is especially striking given the effects of exposure during the experiment. Participants became more informed and more trusting of their chosen outlets and of the press, and trust in the outlet was a strong predictor of outlet choice at baseline, yet consumers were unwilling to pay even heavily discounted prices to sustain that exposure. The result creates a wedge between the effects of professional news consumption and consumers'

private demand for it: temporary subsidies can increase professional news consumption while they are in place without generating sustained private demand once they expire.

6 Heterogeneous effects

Subsidizing news subscriptions impacted trust, political beliefs and attitudes, and news knowledge. We use pre-registered heterogeneous treatment effects to interpret possible mechanisms that could be driving the results. We focus on six outcome variables summarizing the main outcomes of the study: the opinions about the press index, the information index, affective polarization between Lula and Bolsonaro, the political engagement index, the trust in elections index, and the pro-democracy attitudes index.

We pre-registered five variables to explore heterogeneity: news knowledge, news consumption, social media consumption, political position, and outlet choice.²⁷ We split the sample at the median for the first three variables, while political position is categorical.

The estimates point to one key moderator of treatment effects: political ideology. Among participants who identify as centrist or right-wing, treatment effects are generally stronger, with the exception of political engagement (Appendix Table A.27). Specifically, the difference between non-left and left-wing respondents is positive and significant for trust in the chosen newspaper (+0.23 s.d., $p = 0.01$), the index of press opinions (+0.16 s.d., $p = 0.04$), and trust in the electoral system (+0.14 s.d., $p = 0.04$). The difference is negative for the willingness-to-vote index (−0.15 s.d., $p = 0.08$), meaning a higher pro-democratic shift, and the political engagement index (−0.15 s.d., $p = 0.07$). Splitting between centrist and right-wing subsamples reveals similar but noisier estimates (Appendix Table A.26).

Taken together, the results show more positive political spillover effects for the non-left subsample. Trust in the newspaper, media, elections, and pro-democratic attitudes are all higher. This subsample has lower baseline levels of these variables, suggesting its members had more room to update their beliefs.

The overall political ideology pattern is consistent with the broader context of the 2022 elections and the way in which right-wing candidates positioned themselves relative to democracy, the electoral system, and the media. The main presidential candidate, Jair Bolsonaro, was antagonistic toward electoral courts and the media. Direct consumption of professional newspapers can therefore play a role in counteracting top-down speech against media and electoral institutions, even if they are ideologically congruent.

We do not find evidence of a clear role for other moderators, and no single outlet among the five seems to consistently drive results. Tables A.21, A.23, and A.24 report the above-minus-below-median differential in the treatment effect for each index.

²⁷The pre-analysis plan also lists age, gender, and household income as standard moderators. We find little heterogeneity through these moderators (Appendix Tables A.28–A.30).

7 Conclusion

This paper investigates the causal effects of consuming professional newspapers during the 2022 Brazilian general elections. By partnering with five news outlets to provide free three-month digital subscriptions and reading incentives to 1,658 social media users, we offer novel evidence on the political consequences of professional news consumption in a democratic setting.

Our research design led to a 70% take-up of the subscriptions, substantially higher than similar studies promoting news consumption through subscriptions, which allows us to estimate effects of exposure to professional news on outcomes. Furthermore, we provided naturalistic variation by allowing participants to choose among five outlets spanning the political spectrum. In line with theories of selective exposure, participants chose ideologically aligned outlets. The treatment successfully shifted news consumption away from social media and toward the chosen outlet, affecting the participants' news diets.

The findings reveal effects that are important for democracy. First, the treatment increased news knowledge, as measured by news quizzes administered through phone surveys and the endline survey. Second, participants developed more positive perceptions of professional news, reflected in higher trust in their chosen outlet and greater confidence in journalists and the press. Third, we find significant political spillovers: treated participants displayed higher pro-democratic attitudes and political engagement.

Despite the ideologically aligned nature of outlet choices, we find no evidence that the treatment increased affective polarization. This suggests that promoting ideologically aligned but higher-quality professional news consumption does not necessarily lead to increased polarization, even if political engagement increases.

The heterogeneous treatment effects analysis suggests that these results are not driven by a single outlet, but are instead distributed across multiple outlets, supporting the interpretation that the effects are attributable to professional news consumption in general rather than to the editorial line of a specific publication.

Furthermore, the opinion effects are concentrated among centrist and right-wing participants, who started with lower confidence in the press, less trust in elections, and weaker democratic norms, and therefore had more room to update. This pattern is also consistent with professional news serving as a counterpoint to the broader information diet, given the study's context where right-wing candidates were more prone to promote distrust in the electoral system and undemocratic actions.

However, the demand-side findings present a sobering counterpoint. Post-election demand for subscriptions was virtually zero and unaffected by the experiment: two months after the end of the study, only one treated and two control participants subscribed to one of the newspapers. This stands in contrast to results from settings with foreign newspapers (Chen and Yang, 2019), and is plausibly explained by the availability of substitute goods in a free media environment, where social media provides indirect access to professional journalism without requiring a subscription.

Taken together, the results present a policy conundrum: the consumption of professional newspapers can lead to positive political spillovers, but newspapers are scarcely demanded at market prices. Standard policy responses, such as subsidies or public provision of high-quality media, face well-known risks of state capture that could undermine the accountability role of the press (Besley and Prat, 2006). Revenue-sharing arrangements between social media platforms and news organizations could help relax paywalls, but such agreements have not yet materialized at scale. Understanding how to sustain professional journalism in a media environment dominated by social media remains a first-order challenge and requires future research examining how to generate sustained valuation for professional news.

References

- Abadie, A., Athey, S., Imbens, G. W., and Wooldridge, J. M. (2023). When Should You Adjust Standard Errors for Clustering? *The Quarterly Journal of Economics*, 138(1):1–35.
- Allcott, H., Braghieri, L., Eichmeyer, S., and Gentzkow, M. (2020). The welfare effects of social media. *American Economic Review*, 110(3):629–676.
- Anderson, M. L. (2008). Multiple Inference and Gender Differences in the Effects of Early Intervention: A Reevaluation of the Abecedarian, Perry Preschool, and Early Training Projects. *Journal of the American Statistical Association*, 103(484):1481–1495.
- Batista Pereira, F., Bueno, N. S., Nunes, F., and Pavão, N. (2023). Inoculation Reduces Misinformation: Experimental Evidence from Multidimensional Interventions in Brazil. *Journal of Experimental Political Science*, pages 1–12.
- Belloni, A., Chernozhukov, V., and Hansen, C. (2014). Inference on Treatment Effects after Selection among High-Dimensional Controls. *The Review of Economic Studies*, 81(2):608–650.
- Benjamini, Y. and Hochberg, Y. (1995). Controlling the False Discovery Rate: A Practical and Powerful Approach to Multiple Testing. *Journal of the Royal Statistical Society: Series B (Methodological)*, 57(1):289–300.
- Benjamini, Y. and Yekutieli, D. (2001). The control of the false discovery rate in multiple testing under dependency. *The Annals of Statistics*, 29(4).
- Besley, T. and Prat, A. (2006). Handcuffs for the Grabbing Hand? Media Capture and Government Accountability. *American Economic Review*, 96(3):720–736.
- Briole, S., Cagé, J., and Prat, A. (2026). Making Teenagers Read Newspapers. A Nationwide Experiment in French High Schools.
- Broockman, D. and Kalla, J. (2025). Consuming cross-cutting media causes learning and moderates attitudes: A field experiment with Fox News viewers. *The Journal of Politics*.
- Bursztyn, L., Rao, A., Roth, C., and Yanagizawa-Drott, D. (2023). Opinions as Facts. *The Review of Economic Studies*, 90(4):1832–1864.
- Cagé, J. (2020). Media competition, information provision and political participation: Evidence from French local newspapers and elections, 1944–2014. *Journal of Public Economics*, 185.
- Chen, Y. and Yang, D. Y. (2019). The impact of media censorship: 1984 Or brave new world? *American Economic Review*, 109(6):2294–2332.
- Chopra, F., Haaland, I., and Roth, C. (2022). Do People Demand Fact-Checked News? Evidence From U.S. Democrats. *Journal of Public Economics*, 205.

- Chopra, F., Haaland, I., and Roth, C. (2024). The Demand for News: Accuracy Concerns versus Belief Confirmation Motives. *The Economic Journal*, 134(661):1806–1834.
- Council of Europe (2022). Recommendation CM/Rec(2022)4 of the Committee of Ministers to member States on promoting a favourable environment for quality journalism in the digital age. Technical Report CM/Rec(2022)4, Council of Europe.
- DellaVigna, S. and Kaplan, E. (2007). The Fox News effect: media bias and voting. *Quarterly Journal of Economics*, 122(3).
- Einav, L., Klopach, B., and Mahoney, N. (2025). Selling Subscriptions. *American Economic Review*, 115(5).
- Enikolopov, R., Makarin, A., and Petrova, M. (2020). Social Media and Protest Participation: Evidence From Russia. *Econometrica*, 88(4):1479–1514.
- Fithian, W. and Lei, L. (2022). Conditional calibration for false discovery rate control under dependence. *Annals of Statistics*, 50(6):3091–3118.
- Galeotti, A. and Goyal, S. (2010). The Law of the Few. *American Economic Review*, 100(4):1468–1492.
- Gehlbach, S. and Sonin, K. (2014). Government control of the media. *Journal of Public Economics*, 118:163–171.
- Gentzkow, M. (2006). Television and Voter Turnout. *The Quarterly Journal of Economics*, 121(3):931–972.
- Gentzkow, M. and Shapiro, J. M. (2010). What Drives Media Slant? Evidence From U.S. Daily Newspapers. *Econometrica*, 78(1):35–71.
- Gentzkow, M., Shapiro, J. M., and Sinkinson, M. (2011). The Effect of Newspaper Entry and Exit on Electoral Politics. *American Economic Review*, 101(7):2980–3018.
- Gerber, A. S., Karlan, D., and Bergan, D. (2009). Does the media matter? A field experiment measuring the effect of newspapers on voting behavior and political opinions. *American Economic Journal: Applied Economics*, 1(2):35–52.
- Ghanem, D., Hirshleifer, S., and Ortiz-Becerra, K. (2026). Testing attrition bias in field experiments. *Journal of Human Resources*, 61(2):690–731.
- Graham, M. H. and Svobik, M. W. (2020). Democracy in America? Partisanship, Polarization, and the Robustness of Support for Democracy in the United States. *American Political Science Review*, 114(2):392–409.
- Guess, A. M., Malhotra, N., Pan, J., Barberá, P., Allcott, H., Brown, T., Crespo-Tenorio, A., Dimmery, D., Freelon, D., Gentzkow, M., González-Bailón, S., Kennedy, E., Kim, Y. M., Lazer, D., Moehler, D., Nyhan, B., Rivera, C. V., Settle, J., Thomas, D. R., Thorson, E., Tromble, R., Wilkins, A., Wojcieszak, M., Xiong, B., de Jonge, C. K., Franco, A., Mason, W., Stroud, N. J.,

- and Tucker, J. A. (2023a). How do social media feed algorithms affect attitudes and behavior in an election campaign? *Science*, 381(6656):398–404.
- Guess, A. M., Malhotra, N., Pan, J., Barberá, P., Allcott, H., Brown, T., Crespo-Tenorio, A., Dimmery, D., Freelon, D., Gentzkow, M., González-Bailón, S., Kennedy, E., Kim, Y. M., Lazer, D., Moehler, D., Nyhan, B., Rivera, C. V., Settle, J., Thomas, D. R., Thorson, E., Tromble, R., Wilkins, A., Wojcieszak, M., Xiong, B., de Jonge, C. K., Franco, A., Mason, W., Stroud, N. J., and Tucker, J. A. (2023b). Reshares on social media amplify political news but do not detectably affect beliefs or opinions. *Science*, 381(6656):404–408.
- Hopkins, D. J. and Gorton, T. (2024). Unsubscribed and Undemanding: Partisanship and the Minimal Effects of a Field Experiment Encouraging Local News Consumption. *American Journal of Political Science*, 68(4):1217–1233.
- Instituto Brasileiro de Geografia e Estatística (2023). Pesquisa Nacional por Amostra de Domicílios Contínua: Acesso à Internet e à Televisão e Posse de Telefone Móvel Celular para Uso Pessoal 2022. Technical report, Instituto Brasileiro de Geografia e Estatística.
- Kalmoe, N. P., Gubler, J. R., and Wood, D. A. (2018). Toward Conflict or Compromise? How Violent Metaphors Polarize Partisan Issue Attitudes. *Political Communication*, 35(3):333–352.
- Lazarsfeld, P. F., Berelson, B., and Gaudet, H. (1948). *The People’s Choice: How the Voter Makes Up His Mind in a Presidential Campaign*. Columbia University Press, New York.
- Lee, D. S. (2009). Training, Wages, and Sample Selection: Estimating Sharp Bounds on Treatment Effects. *The Review of Economic Studies*, 76(3):1071–1102.
- Levy, R. (2021). Social Media, News Consumption, and Polarization: Evidence from a Field Experiment. *American Economic Review*, 111(3):831–870.
- Martin, G. J. and Yurukoglu, A. (2017). Bias in cable news: Persuasion and polarization. *American Economic Review*, 107(9):2565–2599.
- Moehring, A. and Molina, C. (2026). Social Influence and News Consumption.
- Newman, N., Fletcher, R., Robertson, C. T., Eddy, K., and Nielsen, R. K. (2022). Reuters Institute Digital News Report 2022. Technical report, Reuters Institute for the Study of Journalism.
- Nyhan, B., Settle, J., Thorson, E., Wojcieszak, M., Barberá, P., Chen, A. Y., Allcott, H., Brown, T., Crespo-Tenorio, A., Dimmery, D., Freelon, D., Gentzkow, M., González-Bailón, S., Guess, A. M., Kennedy, E., Kim, Y. M., Lazer, D., Malhotra, N., Moehler, D., Pan, J., Thomas, D. R., Tromble, R., Rivera, C. V., Wilkins, A., Xiong, B., De Jonge, C. K., Franco, A., Mason, W., Stroud, N. J., and Tucker, J. A. (2023). Like-minded sources on Facebook are prevalent but not polarizing. *Nature*, 620(7972):137–144.
- Pew Research Center (2020). Measuring News Consumption in a Digital Era.

- Prior, M. (2009). The Immensely Inflated News Audience: Assessing Bias in Self-Reported News Exposure. *Public Opinion Quarterly*, 73(1):130–143.
- Romano, J. P., Shaikh, A. M., and Wolf, M. (2010). Hypothesis Testing in Econometrics. *Annual Review of Economics*, 2(1):75–104.
- Semenova, V. (2025). Generalized Lee Bounds. *Journal of Econometrics*, 251.
- Stiglitz, J. E., Schiffrin, A., and Groves, D. W. (2022). *Journalism for Development: The role of journalism in promoting democracy, political accountability and sustainable development*. United Nations.
- Szeidl, A. and Szucs, F. (2021). Media Capture Through Favor Exchange. *Econometrica*, 89(1):281–310.
- Trexler, A. (2026). The Minimal Effects of Making Local News Free: Evidence from a Field Experiment. *The Journal of Politics*, 88(3):998–1010.
- Ventura, T., Majumdar, R., and Nagler, J. (2025). Misinformation Beyond Traditional Feeds: Evidence from a WhatsApp Deactivation Experiment in Brazil.
- Vosoughi, S., Roy, D., and Aral, S. (2018). The spread of true and false news online. *Science*, 359(6380):1146–1151.

Figures and Tables






Figure 2: Experimental design and group sizes

Notes: The experimental design, from recruitment through social-media advertisements and baseline screening to random assignment across the control, Subscription (S), and Subscription+Incentives (SI) arms. Nodes report the number of participants at each stage. Section 3 describes recruitment, randomization, and the survey waves in detail.



(A) Daily news consumption over treatment period



(B) Daily news consumption for those browsing, number browsing in a given day

Figure 3: Daily news consumption from outlets' administrative data

Notes: Average daily news consumption at the chosen outlet over the treatment period, measured in the outlets' administrative browsing data. Panel A averages over all treated participants (orange, akin to an intention-to-treat measure) and over treated participants present in the data at least once (blue). Panel B plots the number browsing on a given day (gray, right-hand axis) and average consumption among those who browsed (green, left-hand axis).

Treatment effect (standard deviations)




Figure 4: Treatment effects on news consumption over time, by survey wave

Notes: The SI versus Control difference in three news-diet outcomes – total consumption, websites and apps, and social media – estimated separately by survey wave. Each estimate is $\hat{\beta}_{SI}$ from Equation (1) with post-double selection controls, with outcomes standardized within wave by the Control mean and standard deviation. The baseline wave is pre-treatment and reports randomization balance. Bars are 90%, 95%, and 99% confidence intervals from heteroskedasticity-robust standard errors; sample sizes vary by wave.




Figure 5: Treatment effects on news diet

Notes: Estimates are $\hat{\beta}_{SI}$ from Equation (1), with controls selected separately for each outcome by post-double selection over the baseline variables described in Section 3.7 (Belloni et al., 2014). Standard errors are heteroskedasticity-robust, and bars are 90%, 95%, and 99% confidence intervals. Index outcomes are variance-covariance weighted composites of their family (Anderson, 2008). Sample sizes vary by outcome. “Daily minutes of news consumption” is winsorized at 120 minutes, and “radio” and “print media” are omitted from the main-medium outcomes for negligible response shares. Appendix Tables A.5 and A.6 report these estimates. Subsequent figures follow these conventions unless stated otherwise.




Figure 6: Treatment effects on news and political knowledge

Notes: Estimates are $\hat{\beta}_{SI}$ from Equation (1) for news and political knowledge quiz scores, one estimate per quiz and per index. Quizzes award one point for a correct answer, half a point for “I don’t know,” and zero for an incorrect answer; partial responses are included, and Appendix C lists the statements fielded in each quiz. Appendix Table A.7 reports these estimates. See the note to Figure 5 for conventions.




Figure 7: Treatment effects on opinions about participants' chosen outlet and the press

Notes: Estimates are $\hat{\beta}_{SI}$ from Equation (1) for opinions about the chosen outlet and about the press, in the endline survey sample. The control group was asked about the outlet it chose at baseline even though it received no subscription, the trust-in-other-sources index weights its two components 0.50 each, and the press-opinions index weights its three components 0.43, 0.28, and 0.29. Appendix Table A.8 reports these estimates. See the note to Figure 5 for conventions.




Figure 8: Treatment effects on political engagement

Notes: Estimates are $\hat{\beta}_{SI}$ from Equation (1) for political engagement outcomes, in the endline survey sample. Seven of the eight index components are measured on a five-point scale of the intensity of the action and campaign donation on a four-point scale; the figure also reports following politics (four-point scale) and turnout in each round (indicators), which are not index components. The index weights its eight components 0.11, 0.05, 0.07, 0.09, 0.31, 0.05, 0.19, and 0.12. Appendix Table A.9 reports these estimates. See the note to Figure 5 for conventions.




Figure 9: Treatment effects on affective polarization

Notes: Estimates are $\hat{\beta}_{SI}$ from Equation (1) for affective polarization and affection outcomes, in the endline survey sample. Affection is measured on a 0–100 “feelings thermometer” and shaded from blue to red, and affective polarization is the absolute difference in affection between two political actors or entities. Appendix Table A.10 reports these estimates. See the note to Figure 5 for conventions.




Figure 10: Treatment effects on willingness to vote

Notes: Estimates are $\hat{\beta}_{SI}$ from Equation (1) for willingness to vote for a politician taking each hypothetical undemocratic action, in the endline survey sample. Outcomes keep their natural sign: a higher value means the participant is more willing to vote for the politician, so negative estimates are pro-democratic shifts. The index weights its six components 0.07, 0.19, 0.28, 0.14, 0.05, and 0.28. Appendix Table A.11 reports these estimates. See the note to Figure 5 for conventions.




Figure 11: Treatment effects on trust in the electoral system

Notes: Estimates are $\hat{\beta}_{SI}$ from Equation (1) for trust in the Brazilian electoral system, in the endline survey sample; the three auditable-vote outcomes use only the 92% of respondents aware of the “auditable printed vote” proposal. The index weights its three components 0.44, 0.21, and 0.34. Appendix Table A.12 reports these estimates. See the note to Figure 5 for conventions.




Figure 12: Treatment effects on interest in a post-study subscription

Notes: Estimates are $\hat{\beta}_{SI}$ from Equation (1) for accepting a discounted post-study subscription offer made at endline, an indicator; estimates are shares of participants. Points give the effect for the whole sample, split at the median of baseline news consumption and of baseline trust, and separately by the outlet chosen at baseline. See the note to Figure 5 for conventions.

Table 2: Summary statistics of browsing data

Variable	Mean	p10, p25, p50, p75, p90
<i>Panel A: All newspapers (61% of redeemed coupons in browsing dataset: 322 out of 526)</i>		
Total URLs	34.62	1, 3, 10, 36.75, 86.8
Total browsing duration	10439.87	109.95, 346.62, 1969.65, 7689.67, 28737.25
Total sessions	37.71	1, 3, 12, 32, 116.6
Number of days browsing	15.72	1, 2, 7, 22.75, 46.9
<i>Panel B: Carta Capital (39% of redeemed coupons in browsing dataset: 45 out of 116)</i>		
Total URLs	19.51	1, 2, 4, 10, 22.4
Total browsing duration	3701.25	109.95, 200.1, 534.96, 1278.28, 3807.88
Total sessions	7.42	1, 1, 3, 7, 14.2
Number of days browsing	5.36	1, 1, 2, 5, 11.4
<i>Panel C: Estadão (82% of redeemed coupons in browsing dataset: 121 out of 148)</i>		
Total URLs	45	2, 6, 20, 59, 115
Total browsing duration	16016.23	297.77, 895.8, 4251.11, 18710.63, 43713.7
Total sessions	47.21	1, 4, 17, 57, 124
Number of days browsing	21.17	1, 4, 12, 33, 53
<i>Panel D: Gazeta do Povo (84% of redeemed coupons in browsing dataset: 42 out of 50)</i>		
Total URLs	70.64	4.1, 10, 33, 74.75, 215.1
Total browsing duration	13004.19	389.89, 1198.88, 4249.43, 7882.43, 38186.72
Total sessions	42.81	2.1, 4, 18, 41.75, 119.9
Number of days browsing	19.57	1.1, 3, 11.5, 25.75, 53.2
<i>Panel E: O Globo (35% of redeemed coupons in browsing dataset: 50 out of 142)</i>		
Total URLs	15.68	1, 2.25, 6.5, 22.75, 41.1
Total browsing duration	2730.82	9.8, 35.75, 369.5, 3135.5, 5807.3
Number of days browsing	12.38	1, 2, 5, 19.25, 29.6
<i>Panel F: Valor Econômico (91% of redeemed coupons in browsing dataset: 64 out of 70)</i>		
Total URLs	16.8	1, 2, 7, 16.5, 52.9
Total browsing duration	8975	58.2, 342.5, 1817, 4620, 24043.3
Number of days browsing	12.77	1, 2, 6, 12.75, 40.3

Note: Descriptive statistics on browsing activity in the outlets' administrative data, restricted to participants who redeemed a coupon. Panels report the mean and selected percentiles (10th, 25th, 50th, 75th, 90th) of URLs accessed, browsing duration in seconds, browsing sessions, and distinct days with activity. Percentages in parentheses are the share of that outlet's redeemed coupons that appear in the browsing data. Browsing sessions are not available for O Globo and Valor Econômico, so Panels E and F omit that row.

Table 3: Comparison of news outlet consumption by S and SI groups

	Redeemed	In browsing dataset	Days browsing	Total URLs	Days browsing	Total URLs
	(1)	(2)	(3)	(4)	(5)	(6)
Subscription + Incentives (<i>SI</i>)	0.010 (0.054) [0.85]	0.117 (0.053) [0.02]	−0.564 (3.970) [0.88]	3.498 (9.335) [0.70]	2.185 (1.580) [0.16]	5.147 (3.780) [0.17]
Subscription (<i>S</i>) mean	0.69	0.34	15.85	32.19	5.35	10.87
Subscription (<i>S</i>) std. dev.	0.47	0.48	19.42	41.97	13.45	28.53
Observations	751	751	322	322	751	751

Note: Differences between the Subscription + Incentives (*SI*) and Subscription-only (*S*) groups. Columns 1–2 report coupon redemption and appearance in the browsing data; columns 3–6 report days with activity and URLs accessed, for those observed in the browsing data (columns 3–4) and for the full sample (columns 5–6). Heteroskedasticity-robust standard errors are in parentheses, *p*-values in square brackets. The final rows report the *S*-group mean and standard deviation.

Appendix

A Additional figures and tables




Figure A.1: Example of an ad used for sample recruitment

Notes: On the left, original ads ran in Portuguese. On the right, a version with relevant translations. The full text from the ad read:

- **Portuguese:** “Participe da Pesquisa Acadêmica de Notícias, conduzida por pesquisadores dos departamentos de Economia da Fundação Getúlio Vargas e da Universidade de Stanford. Você poderá ser compensado pelo seu tempo com R\$100 em vale-presente da Americanas ou iFood.
*Atenção: somos um grupo apartidário de pesquisadores, e esse post não é uma propaganda política.”
- **English:** “Participate in the Academic News Research, conducted by researchers from the Economics departments of Fundação Getúlio Vargas and Stanford University. You might be rewarded for your time with a R\$100 gift card of Americanas or iFood.
*Attention: we are a non-partisan group of researchers, and this post is not a political advertisement.”




Figure A.2: Balance tests for key baseline variables and samples

Notes: Each point is the standardized difference in means between the *Subscription + Incentives* (SI) and control (C) groups for one baseline covariate, estimated separately within the baseline, endline, pre-elections, and September 7 samples. Vertical lines are 95% confidence intervals from standard errors clustered at the stratification level. Dashed lines mark ± 0.1 standard deviations, and the legend reports each sample's joint significance test p -value.




Figure A.3: Treatment effects on news diet, for participants who thought we were not promoting a specific view on news

Notes: Treatment effects on news-diet outcomes, restricting the sample to participants who did not believe we were promoting a specific view of the news. Estimates are $\hat{\beta}_{SI}$ from Equation (1) with post-double selection controls, as in Figure 5. Sample sizes vary by outcome. Bars are 90%, 95%, and 99% confidence intervals from heteroskedasticity-robust standard errors.




Figure A.4: Bounds for treatment effects

Notes: Treatment effects on the 43 main outcomes, each estimated from Equation (1). Bars are 95% confidence intervals from heteroskedasticity-robust standard errors; gray rectangles are basic Lee bounds, trimming the control group so that response rates are equal across arms, with item non-response counted as attrition. Colored rectangles are the generalized Lee bounds of Semenova (2025), computed for the continuous outcomes only. Sample sizes vary by outcome. See Appendix Section D.5 for how the bounds are constructed.

Table A.1: Summary statistics and balance tests

Variable	(1)	(2)	(3)	(4)	(5)	(6)	(7)
	Summary statistics		Balance tests				
	Mean	p10, p25, p50, p75, p90	Control	Subscription (S)	Subscription + Incentives (SI)	p-val: (4)-(3)	p-val: (5)-(3)
<i>A. Demographics</i>							
Age	44.33	26, 34.25, 44, 54, 62	44.054	44.052	44.721	0.999	0.29
Female	0.73		0.734	0.675	0.731	0.31	0.889
High school	0.23		0.233	0.234	0.221	0.977	0.487
Undergraduate	0.62		0.623	0.675	0.604	0.294	0.377
Graduate	0.15		0.144	0.091	0.175	0.11	0.032
Below R\$2,000	0.23		0.225	0.234	0.231	0.86	0.769
R\$2,000 - R\$5,000	0.35		0.348	0.364	0.358	0.791	0.715
Over R\$5,000	0.42		0.427	0.403	0.411	0.677	0.529
Black	0.11		0.121	0.156	0.092	0.368	0.063
Brown	0.29		0.299	0.286	0.286	0.817	0.595
White/Caucasian	0.58		0.569	0.545	0.598	0.673	0.26
<i>B. News and social media</i>							
News consumption (min.)	61.81	20, 30, 60, 100, 120	62.037	63.078	61.369	0.762	0.708
News knowledge	4.54	3, 4, 4.5, 5.5, 6	4.558	4.383	4.536	0.207	0.672
Social media usage (min.)	115.51	30, 60, 100, 180, 240	116.639	117.584	113.752	0.898	0.371
<i>C. Declared political position</i>							
Left-wing	0.54		0.54	0.597	0.531	0.186	0.648
Right-wing	0.16		0.155	0.182	0.172	0.446	0.314
Center	0.3		0.304	0.221	0.297	0.037	0.624
<i>D. Choice of news outlet</i>							
Choice: Carta Capital	0.23		0.239	0.208	0.228	0.248	0.511
Choice: Estadão	0.26		0.26	0.325	0.26	0.085	0.97
Choice: Gazeta do Povo	0.09		0.095	0.104	0.091	0.727	0.76
Choice: O Globo	0.26		0.255	0.221	0.273	0.327	0.374
Choice: Valor Econômico	0.15		0.151	0.143	0.148	0.823	0.839
Joint significance test						0.88	0.56
Number of observations		1,658	907	77	674		

Note: Baseline characteristics of all randomized participants, one observation per participant. Columns 1–2 report the pooled mean and the 10th, 25th, 50th, 75th, and 90th percentiles; columns 3–5 report means for the control, Subscription (S), and Subscription+Incentives (SI) arms; columns 6–7 report p -values from tests of equal means against control. The penultimate row reports a joint Wald test per comparison, with group sizes in the final row.

Table A.2: Estimates of relationship between choice of outlet and trust and ideological distance

	<i>Dependent variable:</i>			
	choice			
	(1)	(2)	(3)	(4)
Absolute ideological distance = 1				−0.301 (0.113) [<0.01]
Absolute ideological distance = 2				−0.554 (0.123) [<0.01]
Absolute ideological distance ≥ 3				−0.890 (0.130) [<0.01]
Absolute ideological distance = Unknown				−0.933 (0.166) [<0.01]
Trust = Low	1.409 (0.257) [<0.01]		1.267 (0.259) [<0.01]	1.353 (0.258) [<0.01]
Trust = Medium	3.032 (0.254) [<0.01]		2.702 (0.259) [<0.01]	2.814 (0.257) [<0.01]
Trust = High	5.041 (0.268) [<0.01]		4.531 (0.276) [<0.01]	4.650 (0.274) [<0.01]
Trust = Unknown	1.191 (0.290) [<0.01]		1.125 (0.304) [<0.01]	1.232 (0.301) [<0.01]
Absolute ideological distance (linear)		−0.603 (0.030) [<0.01]	−0.249 (0.035) [<0.01]	
Absolute ideological distance unknown			248.010 (34.722) [<0.01]	
Observations	1,658	1,658	1,658	1,658
Log Likelihood	−1,862.898	−2,257.210	−1,831.583	−1,832.375

Note: Multinomial logit estimates of the outlet chosen at baseline, one observation per participant. Regressors are pre-treatment trust in the outlet and the absolute gap between own and perceived outlet ideology, entered linearly in columns 1–3 and categorically in column 4. Asymptotic standard errors are in parentheses, p -values in square brackets. The only terms not shown are the alternative-specific intercepts, one per non-reference outlet, which are nuisance parameters absorbing each outlet’s baseline popularity. The “distance unknown” indicator is subject to separation, so its magnitude is not interpretable; its column-2 estimate is included in the model but not shown.

Table A.3: Coverage of the browsing data by outlet

Outlet	Total URLs accessed	Share missing	Total home page URLs	Total unique URLs accessed	Share of unique URLs missing
Estadão	1559	0.01	1448	90	0.00
Carta Capital	1577	0.93	1430	113	0.11
Gazeta do Povo	380	0.05	274	96	0.04
O Globo	2109	0.06	1913	128	0.10
Valor Econômico	2187	0.03	1884	116	0.01
Total (unweighted)	7812	0.217	6949	543	0.055
Total (weighted)	8473.9	0.225	7624.8	548.9	0.061

Note: Audit of the browsing data, one row per partner outlet. The authors browsed about 100 unique URLs from each outlet on a computer over one week and asked each newspaper to extract those records; the Chrome browsing history is the benchmark. Columns report the URLs accessed in the audit, the share missing from the outlet’s extract, home-page URLs, unique URLs accessed, and the share of unique URLs missing. Carta Capital does not record home-page accesses, so its share of missing URLs is high. “Total (unweighted)” pools the five outlets; “Total (weighted)” reweights outlets by the number of participants who chose each. The table is descriptive and reports no inference; see Appendix C.1.

Table A.4: Beliefs about the experimenter's intentions, by treatment arm

	View of news outlets	Political side
	(1)	(2)
Subscription (<i>S</i>)	−0.013 (0.066) [0.84]	−0.048 (0.066) [0.47]
Subscription + Incentives (<i>SI</i>)	−0.082 (0.026) [<0.01]	0.030 (0.026) [0.24]
Control mean	0.72	0.7
Control std. dev.	0.45	0.46
Observations	1,363	1,363

Note: Beliefs about the experiment's purpose, elicited in the endline survey sample, regressed on indicators for the Subscription (*S*) and Subscription+Incentives (*SI*) arms, control omitted. The outcome is an indicator that the respondent believed the experiment did *not* aim to promote a particular view of news outlets (column 1) or a political side (column 2); coefficients are differences in shares. Heteroskedasticity-robust standard errors are in parentheses, *p*-values in square brackets.

Table A.5: Treatment effects – news diet (Part 1: overall consumption and main medium)

	Overall consumption	Follows politics	Main medium of consumption		
	(minutes, std.)	(4-point scale)	App/sites of outlets	Television	Social media
	(1)	(2)	(3)	(4)	(5)
Subscription (<i>S</i>)	−0.031 (0.119) [0.80]	0.079 (0.105) [0.46]	0.198 (0.120) [0.10]	−0.099 (0.096) [0.30]	−0.087 (0.106) [0.41]
Subscription + Incentives (<i>SI</i>)	0.007 (0.049) [0.88]	0.100 (0.048) [0.04]	0.067 (0.053) [0.21]	0.029 (0.050) [0.56]	−0.100 (0.050) [0.05]
Control raw mean	70.37	1.99	0.42	0.25	0.29
Control raw std. dev.	41.72	0.82	0.49	0.43	0.46
Observations	1,363	1,363	1,363	1,363	1,363

Note: Estimates are $\hat{\beta}_S$ and $\hat{\beta}_{SI}$ from Equation (1), with controls selected separately for each outcome by post-double selection over a set of baseline variables (see Section 3.7). Standard errors are heteroskedasticity-robust, reported in parentheses, with p -values in square brackets. Index outcomes are variance-covariance weighted composites of their family (Anderson, 2008). Sample sizes vary by outcome and are reported in the final row. “Daily minutes of news consumption” is winsorized at 120 minutes, and “radio” and “print media” are omitted from the main-medium outcomes for negligible response shares. Subsequent tables follow these conventions unless stated otherwise.

Table A.6: Treatment effects – news diet (Part 2: intensity of each medium)

	Intensity of each medium of consumption (5-point scale)						
	App/sites of outlets	Social media	Cable TV	Local TV	National TV	Video apps	Radio/podcasts
	(1)	(2)	(3)	(4)	(5)	(6)	(7)
Subscription (<i>S</i>)	0.061 (0.118) [0.60]	−0.102 (0.108) [0.35]	−0.058 (0.094) [0.54]	−0.059 (0.089) [0.50]	−0.073 (0.068) [0.28]	0.174 (0.093) [0.06]	0.058 (0.125) [0.64]
Subscription + Incentives (<i>SI</i>)	0.202 (0.047) [<0.01]	−0.082 (0.048) [0.09]	0.023 (0.043) [0.60]	0.072 (0.042) [0.08]	0.049 (0.042) [0.24]	−0.019 (0.044) [0.67]	0.003 (0.051) [0.96]
Control raw mean	1.81	2.38	1.35	1.66	1.93	1.39	1.28
Control raw std. dev.	1.22	1.16	1.35	1.31	1.29	1.29	1.16
Observations	1,363	1,363	1,363	1,363	1,363	1,363	1,363

Note: Estimates are $\hat{\beta}_S$ and $\hat{\beta}_{SI}$ from Equation (1) for the intensity of news consumption on each medium, each measured on a five-point scale, in the endline survey sample. The same estimates for the *SI* group appear in Figure 5. See the note to Table A.5 for conventions.

Table A.7: Treatment effects – news knowledge and political knowledge

	Overall score	Correct news score	Incorrect news score	Overall score	Correct news score	Incorrect news score	Fake news score	Overall score	Correct news score	Incorrect news score	Fake news score	Index of news knowledge	Overall score
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	(10)	(11)	(12)	(13)
Subscription (<i>S</i>)	0.019 (0.104) [0.86]	−0.035 (0.109) [0.75]	−0.017 (0.115) [0.88]	0.156 (0.121) [0.20]	0.144 (0.159) [0.36]	−0.088 (0.114) [0.44]	0.118 (0.147) [0.42]	0.221 (0.110) [0.04]	0.049 (0.118) [0.68]	0.205 (0.118) [0.08]	0.177 (0.114) [0.12]	0.238 (0.107) [0.03]	−0.219 (0.120) [0.07]
Subscription + Incentives (<i>SI</i>)	0.069 (0.054) [0.20]	0.021 (0.060) [0.73]	0.096 (0.054) [0.08]	0.126 (0.057) [0.03]	0.222 (0.062) [<0.01]	−0.017 (0.057) [0.76]	−0.012 (0.063) [0.85]	0.044 (0.049) [0.37]	0.105 (0.053) [0.05]	−0.006 (0.052) [0.91]	−0.011 (0.053) [0.84]	0.100 (0.053) [0.06]	−0.009 (0.047) [0.85]
Control raw mean	3.98	2.77	1.27	5.4	2.13	1.29	2.04	6.39	2.78	1.87	1.74	0	5.35
Control raw std. dev.	1.14	0.76	0.62	1.33	0.64	0.82	0.69	1.53	0.75	0.75	0.81	1.02	1.45
Observations	1,320	1,242	1,291	1,207	1,179	1,177	1,192	1,363	1,363	1,363	1,363	1,088	1,363

Note: Estimates are $\hat{\beta}_S$ and $\hat{\beta}_{SI}$ from Equation (1) for news and political knowledge quiz scores. Columns 1 to 3 use the September 7 sample, columns 4 to 7 the pre-elections sample, and columns 8 to 11 and 13 the endline sample; column 12 indexes the scores in columns 1, 4, and 8, with weights 0.37, 0.37, and 0.26. See the note to Table A.5 for conventions.

Table A.8: Treatment effects – opinions about participants’ chosen outlet and the press

	Trust in chosen outlet (1)	Trust in chosen outlet, medium or high at baseline (2)	Trust in chosen outlet, low or undeclared at baseline (3)	Trust in friends and family news (4)	Trust in social media news (5)	Index of trust in other sources (6)	Updated their perceived ideology (7)	Updated their perceived ideology, undeclared at baseline (8)	Confidence in press (9)	Journalists provide fair coverage of elections (10)	Journalists provide fair coverage of politics (11)	Index of press opinions (12)
Subscription (<i>S</i>)	0.219 (0.110) [0.05]	0.133 (0.126) [0.29]	0.759 (0.260) [<0.01]	0.022 (0.120) [0.85]	−0.198 (0.118) [0.09]	−0.125 (0.115) [0.28]	−0.299 (0.174) [0.09]	0.559 (0.099) [<0.01]	0.036 (0.104) [0.72]	0.002 (0.093) [0.98]	−0.047 (0.102) [0.64]	0.017 (0.095) [0.86]
Subscription + Incentives (<i>SI</i>)	0.300 (0.046) [<0.01]	0.234 (0.047) [<0.01]	0.714 (0.174) [<0.01]	−0.078 (0.047) [0.10]	−0.045 (0.051) [0.38]	−0.079 (0.048) [0.10]	−0.001 (0.071) [0.99]	0.281 (0.139) [0.05]	0.038 (0.041) [0.36]	0.089 (0.044) [0.05]	0.105 (0.044) [0.02]	0.088 (0.038) [0.02]
Control raw mean	2.42	0.15	−1.02	1.39	1.73	0	0.51	0.42	1.63	1.79	1.66	0
Control raw std. dev.	0.98	0.9	1.02	0.87	0.8	1.01	0.5	0.5	0.76	0.69	0.7	1
Observations	1,363	1,184	179	1,363	1,363	1,363	819	91	1,363	1,363	1,363	1,363

Note: Estimates are $\hat{\beta}_S$ and $\hat{\beta}_{SI}$ from Equation (1) for opinions about the chosen outlet and about the press, in the endline survey sample. Column 6 is the index of columns 4 and 5, with weights 0.50 each; column 12 is the index of columns 9 to 11, with weights 0.43, 0.28, and 0.29. The same estimates for the *SI* group appear in Figure 7. See the note to Table A.5 for conventions.

Table A.9: Treatment effects – political engagement

	Talks about politics with friends and family	Shares political content on WhatsApp/Telegram	Shares political content on social media	Volunteers in electoral campaign	Donates money for campaigns	Persuades others to vote for a candidate	Watches debates for governor	Reads interviews of candidates	Index of political engagement	Turnout (round 1)	Turnout (round 2)	Follows politics
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	(10)	(11)	(12)
Subscription (<i>S</i>)	0.117 (0.076) [0.13]	0.201 (0.093) [0.03]	0.121 (0.076) [0.11]	0.100 (0.117) [0.39]	0.121 (0.155) [0.44]	0.127 (0.116) [0.27]	0.316 (0.110) [<0.01]	0.088 (0.101) [0.38]	0.247 (0.122) [0.04]	−0.163 (0.141) [0.25]	−0.147 (0.136) [0.28]	0.079 (0.105) [0.46]
Subscription + Incentives (<i>SI</i>)	−0.026 (0.044) [0.55]	−0.030 (0.040) [0.46]	0.032 (0.040) [0.42]	0.035 (0.043) [0.42]	0.021 (0.052) [0.69]	0.075 (0.044) [0.09]	0.085 (0.049) [0.09]	0.073 (0.049) [0.14]	0.065 (0.040) [0.11]	0.032 (0.052) [0.54]	0.068 (0.051) [0.18]	0.100 (0.048) [0.04]
Control raw mean	2.38	1.4	1.41	0.76	0.1	1.34	1.88	2.43	0	0.92	0.91	1.99
Control raw std. dev.	0.85	1.15	1.21	1.15	0.42	1.29	1.31	0.98	0.99	0.28	0.29	0.82
Observations	1,363	1,363	1,363	1,363	1,363	1,363	1,363	1,363	1,363	1,363	1,363	1,363

Note: Estimates are $\hat{\beta}_S$ and $\hat{\beta}_{SI}$ from Equation (1) for political engagement outcomes, all in the endline survey sample. Column 9 is the index of the scores in columns 1 to 8, with weights 0.11, 0.05, 0.07, 0.09, 0.31, 0.05, 0.19, and 0.12. The same estimates for the *SI* group appear in Figure 8. See the note to Table A.5 for conventions.

Table A.10: Treatment effects – affective polarization

	Bolsonaro x Lula	Bolsonaro x Supreme Court	Left x Right (politicians)	Left x Right (supporters)	Bolsonaro	Lula	Supreme Court	Left-wing politicians	Right-wing politicians	Left-wing supporters	Right-wing supporters
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	(10)	(11)
Subscription (<i>S</i>)	0.061 (0.091) [0.50]	0.160 (0.089) [0.07]	−0.060 (0.103) [0.56]	0.023 (0.097) [0.81]	−0.155 (0.060) [<0.01]	0.079 (0.084) [0.35]	0.185 (0.083) [0.03]	0.007 (0.094) [0.94]	0.005 (0.067) [0.94]	0.037 (0.092) [0.69]	−0.084 (0.080) [0.30]
Subscription + Incentives (<i>SI</i>)	−0.004 (0.042) [0.92]	0.018 (0.043) [0.68]	0.005 (0.043) [0.91]	−0.006 (0.043) [0.88]	0.016 (0.033) [0.63]	0.016 (0.032) [0.62]	0.055 (0.035) [0.11]	−0.043 (0.035) [0.22]	0.026 (0.038) [0.49]	−0.027 (0.036) [0.46]	0.060 (0.040) [0.13]
Control raw mean	70.04	65.84	57.98	59.2	21.51	64.57	63.23	61.63	28.58	59.43	24.37
Control raw std. dev.	32.51	31.15	32.54	32.55	33.63	37.08	33.12	33.35	31.05	33.54	32.08
Observations	1,363	1,363	1,363	1,363	1,363	1,363	1,363	1,363	1,363	1,363	1,363

Note: Estimates are $\hat{\beta}_S$ and $\hat{\beta}_{SI}$ from Equation (1) for affective polarization and affection outcomes, all in the endline survey sample. The same estimates for the *SI* group appear in Figure 9. See the note to Table A.5 for conventions.

Table A.11: Treatment effects – willingness to vote

	Contest electoral results in courts if loses by a small margin	Rules by executive orders	Deny authorization for opposition protests	Sue journalists who do not disclose their sources	Disrespect judicial orders	Suggests a temporary military intervention	Willingness-to- vote index
	(1)	(2)	(3)	(4)	(5)	(6)	(7)
Subscription (<i>S</i>)	−0.041 (0.073) [0.57]	−0.082 (0.124) [0.51]	−0.102 (0.128) [0.42]	−0.080 (0.121) [0.51]	0.003 (0.122) [0.98]	0.162 (0.138) [0.24]	−0.026 (0.108) [0.81]
Subscription + Incentives (<i>SI</i>)	−0.025 (0.035) [0.48]	−0.093 (0.048) [0.05]	−0.042 (0.051) [0.41]	−0.002 (0.047) [0.96]	0.068 (0.047) [0.15]	−0.042 (0.050) [0.41]	−0.065 (0.042) [0.13]
Control raw mean	0.83	1.05	0.8	1.03	0.85	1.21	0
Control raw std. dev.	1.14	1.08	1.01	1.07	0.98	1.54	1.02
Observations	1,363	1,363	1,363	1,363	1,363	1,363	1,363

Note: Estimates are $\hat{\beta}_S$ and $\hat{\beta}_{SI}$ from Equation (1) for willingness to vote for a politician taking each undemocratic action, in the endline survey sample. Outcomes keep their natural sign: a higher value means the participant is more willing to vote for the politician, so negative estimates are pro-democratic shifts. Column 7 is the willingness-to-vote index of columns 1–6, with weights 0.07, 0.19, 0.28, 0.14, 0.05, and 0.28. The same estimates appear in Figure 10. See the note to Table A.5 for conventions.

Table A.12: Treatment effects – trust in the electoral system

	Agrees elections are free and fair	Trust in electronic voting system	Confidence in elections	Index of trust in electoral system	Aware of the auditable vote	In favor of keeping electronic ballots	In favor of reverting to paper ballots	In favor of the auditable vote proposal
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Subscription (<i>S</i>)	0.049 (0.120) [0.68]	−0.100 (0.079) [0.21]	0.008 (0.076) [0.92]	0.002 (0.087) [0.98]	−0.017 (0.134) [0.90]	0.078 (0.080) [0.33]	−0.176 (0.049) [<0.01]	−0.028 (0.085) [0.74]
Subscription + Incentives (<i>SI</i>)	−0.002 (0.044) [0.96]	0.002 (0.031) [0.96]	0.024 (0.033) [0.47]	0.007 (0.032) [0.82]	−0.026 (0.048) [0.59]	0.092 (0.040) [0.02]	−0.034 (0.054) [0.54]	−0.090 (0.043) [0.04]
Control raw mean	0.74	2.31	2.39	0	0.92	0.71	0.02	0.27
Control raw std. dev.	0.44	1.02	1	1.01	0.27	0.45	0.15	0.44
Observations	1,363	1,363	1,363	1,363	1,363	1,258	1,258	1,258

Note: Estimates are $\hat{\beta}_S$ and $\hat{\beta}_{SI}$ from Equation (1) for trust in the Brazilian electoral system. Columns 1–5 use the endline survey sample and columns 6–8 the subsample aware of the “auditable printed vote” proposal. Column 4 is the index of columns 1–3, with weights 0.44, 0.21, and 0.34. The same estimates appear in Figure 11. See the note to Table A.5 for conventions.

Table A.13: Treatment effects – secondary political outcomes

	Voted Bolsonaro (2nd round)	Voted Lula (2nd round)	Confidence in government	Confidence in Congress	Confidence in courts	Support for military rule	Importance of democracy	How democratic Brazil is today
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Subscription (<i>S</i>)	0.001 (0.035) [0.98]	−0.028 (0.057) [0.63]	0.007 (0.103) [0.95]	−0.108 (0.103) [0.30]	0.00003 (0.083) [1.00]	−0.058 (0.061) [0.34]	−0.072 (0.120) [0.55]	0.075 (0.117) [0.52]
Subscription + Incentives (<i>SI</i>)	0.029 (0.029) [0.32]	−0.034 (0.029) [0.24]	−0.092 (0.039) [0.02]	0.025 (0.049) [0.61]	0.036 (0.036) [0.31]	−0.015 (0.033) [0.65]	0.019 (0.046) [0.69]	0.018 (0.046) [0.70]
Control raw mean	0.21	0.73	0.87	1.17	1.83	0.42	9.56	4.8
Control raw std. dev.	0.41	0.44	1.02	0.8	0.98	0.83	1.31	3.03
Scale	0–1	0–1	0–3	0–3	0–3	0–3	0–10	0–10
Observations	1,247	1,247	1,363	1,363	1,363	1,279	1,363	1,363

Note: Estimates are $\hat{\beta}_S$ and $\hat{\beta}_{SI}$ from Equation (1) for the pre-registered secondary political outcomes not reported in the tables above, all in the endline survey sample. Columns 1–2 are indicators for the self-reported second-round vote; columns 3–5 are confidence in government, Congress, and courts (confidence in elections appears in Table A.12); columns 6–8 follow the World Values Survey wording for support for military rule, the importance of democracy, and how democratically Brazil is ruled today. All outcomes are standardized on the control group; the “Scale” row gives each item’s range, and the control mean shows how close to a bound the item sits (importance of democracy averages 9.6 on a 0–10 scale). “Don’t know” answers to the military-rule item are treated as missing. See the note to Table A.5 for conventions.

Table A.14: Treatment effects – family indices with the treatment arms pooled

	Index of news diet (1)	Index of news knowledge (2)	Index of press opinions (3)	Index of trust in other sources (4)	Index of political engagement (5)	Willingness-to-vote index (6)	Index of trust in electoral system (7)
Pooled treatment (<i>S</i> or <i>SI</i>)	0.131 (0.043) [<0.01]	0.108 (0.051) [0.04]	0.081 (0.037) [0.03]	−0.084 (0.047) [0.07]	0.084 (0.040) [0.04]	−0.061 (0.041) [0.14]	0.007 (0.031) [0.83]
Two-arm <i>SI</i> estimate, Equation (1)	0.150 (0.045) [<0.01]	0.100 (0.053) [0.06]	0.088 (0.038) [0.02]	−0.079 (0.048) [0.10]	0.065 (0.040) [0.11]	−0.065 (0.042) [0.13]	0.007 (0.032) [0.82]
Observations	1,363	1,088	1,363	1,363	1,363	1,363	1,363

Note: The coefficient row replaces the two arm indicators of Equation (1) with a single indicator equal to one for participants assigned to either *S* or *SI* (control omitted). Controls are selected by post-double selection with the pooled indicator as the target. The reference rows repeat $\hat{\beta}_{SI}$ from Equation (1), its standard error, and its *p*-value, as reported in Tables A.5–A.12. All estimates use the endline survey sample. Standard errors are heteroskedasticity-robust

Table A.15: Survey completion rates by treatment status and coupon redemption

	Survey sample					
	September 7		Pre-elections survey		Endline survey	
	(1)	(2)	(3)	(4)	(5)	(6)
Subscription (<i>S</i>)	−0.142 (0.051) [<0.01]		−0.154 (0.055) [<0.01]		−0.123 (0.048) [0.01]	
Subscription + Incentives (<i>SI</i>)	−0.133 (0.023) [<0.01]		−0.136 (0.024) [<0.01]		−0.120 (0.020) [<0.01]	
<i>S</i> , used coupon		−0.026 (0.047) [0.57]		−0.036 (0.063) [0.57]		−0.009 (0.049) [0.86]
<i>S</i> , did not use coupon		−0.398 (0.100) [<0.01]		−0.416 (0.084) [<0.01]		−0.377 (0.096) [<0.01]
<i>SI</i> , used coupon		−0.034 (0.020) [0.09]		−0.038 (0.022) [0.09]		−0.022 (0.019) [0.24]
<i>SI</i> , did not use coupon		−0.364 (0.039) [<0.01]		−0.368 (0.046) [<0.01]		−0.349 (0.039) [<0.01]
Control completion rate	0.857	0.857	0.791	0.791	0.877	0.877
Observations	1,658	1,658	1,658	1,658	1,658	1,658

Note: The dependent variable is an indicator for responding to the survey in the column header, over the full randomized sample; columns 1–2, 3–4, and 5–6 use the September 7, pre-elections, and endline surveys. Odd columns use assignment alone, even columns split it by coupon redemption. No controls are included. Standard errors are clustered at the stratification level, with *p*-values in square brackets; see Appendix D.

Table A.16: Outcome-by-outcome tests: balance at baseline and conditional on response (IV-R) in the post-baseline samples

(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	(10)
Family	Outcomes at baseline survey	Baseline sample (Balance: $\pi_{10} = 0$)		Post-baseline surveys samples (IV-R: $\pi_{10} = \pi_{11} = 0$)					
		p	q	Endline		Pre-elections		Sep. 7	
				p	q	p	q	p	q
News diet	1. News consumption (minutes)	0.45	0.71	0.67	0.78	0.81	0.96	0.58	0.87
	2. Closely follows politics	0.2	0.60	0.09*	0.48	0.47	0.94	0.87	0.95
	3. Main medium: App/Sites	0.91	0.98	0.34	0.65	0.81	0.96	0.85	0.95
	4. Main medium: TV	0.95	0.98	0.54	0.70	0.86	0.96	0.14	0.79
	5. Main medium: Social media	0.68	0.93	0.64	0.77	0.76	0.96	0.44	0.86
	6. Intensity: Web news	0.93	0.98	0.28	0.65	0.56	0.94	0.2	0.79
	7. Intensity: Social media news	0.3	0.61	0.71	0.79	0.54	0.94	0.45	0.86
	8. Intensity: Cable TV news	0.46	0.71	0.03*	0.48	0.55	0.94	0.49	0.86
	9. Intensity: Local TV news	0.12	0.50	0.05*	0.48	0.06*	0.88	0.31	0.81
	10. Intensity: National TV news	0.04*	0.39	0.06*	0.48	0.05*	0.88	0.17	0.79
	11. Intensity: Video news	0.69	0.93	0.52	0.70	0.35	0.94	0.95	0.98
	12. Intensity: Radio news	0.83	0.98	0.41	0.65	0.78	0.96	0.87	0.95
	13. Index: News diet	0.87	0.98	0.42	0.65	0.67	0.94	0.98	0.98
News knowledge	14. News knowledge	0.71	0.94	0.14	0.52	0.48	0.94	0.06*	0.79
Opinions about press	15. Trust in chosen newspaper	0.11	0.50	0.35	0.65	0.1	0.94	0.16	0.79
	16. Confidence in press	0.32	0.61	0.15	0.52	0.57	0.94	0.71	0.94
	17. Fair media coverage of elections	0.61	0.88	0.91	0.94	0.96	0.96	0.7	0.94
	18. Fair media coverage of politics	0.13	0.50	0.28	0.65	0.22	0.94	0.21	0.79
	19. Index: press opinions	0.22	0.60	0.45	0.68	0.66	0.94	0.74	0.94
Political engagement	20. Talk about politics	0.99	0.99	0.53	0.70	0.89	0.96	0.72	0.94
	21. Share political content on WhatsApp	0.01*	0.27	0.02*	0.48	0.04*	0.88	0.18	0.79
	22. Share political content on social media	0.1	0.50	0.42	0.65	0.39	0.94	0.25	0.81
	23. Volunteers for campaigns	0.05*	0.39	0.32	0.65	0.32	0.94	0.22	0.79
	24. Donates to campaigns	0.55	0.83	0.7	0.79	0.76	0.96	0.79	0.95
	25. Tries to convince others to vote	0.31	0.61	0.41	0.65	0.57	0.94	0.55	0.86
	26. Watches governor debates	0.94	0.98	0.82	0.89	0.95	0.96	0.37	0.86
	27. Reads interviews with candidates	0.96	0.98	0.85	0.90	0.63	0.94	0.45	0.86
	28. Index: political engagement	0.39	0.65	0.55	0.71	0.58	0.94	0.38	0.86
Willingness-to-vote	29. Accepts election outcome	0.95	0.98	0.97	0.97	0.93	0.96	0.88	0.95
	30. Rule by executive orders	0.38	0.65	0.24	0.64	0.34	0.94	0.49	0.86
	31. Ban protests	0.17	0.58	0.24	0.64	0.24	0.94	0.3	0.81
	32. Prosecute journalists	0.92	0.98	0.08*	0.48	0.94	0.96	0.91	0.96
	33. Ignore court rulings	0.34	0.63	0.6	0.74	0.66	0.94	0.71	0.94
	34. Index: willingness-to-vote	0.31	0.61	0.23	0.64	0.43	0.94	0.55	0.86
Trust in electoral system	35. Elections are free and fair	0.18	0.58	0.1	0.48	0.43	0.94	0.27	0.81
	36. Trust in electronic voting system	0.3	0.61	0.47	0.69	0.88	0.96	0.83	0.95
	37. Confidence in elections	0.01*	0.27	0.09*	0.48	0.26	0.94	0.07*	0.79
	38. Index: trust in electoral system	0.05*	0.39	0.06*	0.48	0.37	0.94	0.17	0.79
Polarization	39. Bolsonaro x Lula	0.06*	0.39	0.23	0.64	0.45	0.94	0.22	0.79
	40. Bolsonaro x Supreme Court	0.27	0.61	0.39	0.65	0.57	0.94	0.4	0.86
	41. Left x Right politicians	0.24	0.61	0.41	0.65	0.65	0.94	0.54	0.86
	42. Left x Right supporters	0.08*	0.47	0.13	0.52	0.41	0.94	0.21	0.79

Note: One row per baseline outcome, grouped by outcome family. Each p comes from a separate regression of the baseline outcome on all 1,658 randomized participants as in Equation (D.5), which compares the SI arm with the control group; the estimated model also includes the S -arm indicator and its interaction with R_i as unrestricted nuisance terms, which the printed equation omits. The baseline-sample column is a balance test: every participant answered the baseline survey ($R_i = 1$), so it tests $\pi_{10} = 0$ alone. The post-baseline columns test $\pi_{10} = \pi_{11} = 0$, the implication of independence between treatment and baseline outcomes conditional on response status (IV-R) that identifies the treatment effect among respondents (Ghanem et al., 2026). Each q is the Benjamini and Hochberg (1995) adjusted p -value, computed within the sample across the 42 outcomes. Standard errors are heteroskedasticity-robust. *: $p < 0.10$.

Table A.17: Complier balance tests – demographic and media covariates

Demographic variable	Δ <i>Complier</i>	Std. Error	Test p -val (= 0)
Age	−0.29	0.07	<0.01*
Female	−0.05	0.08	0.48
High school or less	−0.18	0.08	0.03*
Undergraduate	0.02	0.08	0.79
Graduate	0.19	0.09	0.04*
Income: below R\$2,000	−0.21	0.08	0.01*
Income: R\$2,000-R\$5,000	0.07	0.08	0.32
Income: over R\$5,000	0.11	0.09	0.22
Black	0.06	0.07	0.37
Brown	0.01	0.08	0.94
White	−0.03	0.08	0.66
Left-wing	0.13	0.09	0.15
Centrist	−0.14	0.11	0.19
Right-wing	−0.03	0.08	0.74
Outlet: Carta Capital	−0.05	0.08	0.59
Outlet: Estadão	0.11	0.08	0.16
Outlet: Gazeta do Povo	0.04	0.07	0.61
Outlet: O Globo	0.02	0.10	0.86
Outlet: Valor Econômico	−0.14	0.10	0.16
Social media time	0.13	0.08	0.13
Joint significance test			<0.01

Note: Differences in mean baseline demographics between participants who redeemed their coupon (*compliers*) and those who did not, within the 751 participants assigned to either treatment arm (*S* and *SI* pooled; 526 compliers), one row per covariate; covariates are standardized. The comparison is compliers against non-compliers, not *SI* against control. Column 1 reports the difference, from a regression of the covariate on a complier indicator; column 2 its standard error; and column 3 the p -value, from a t distribution with 749 degrees of freedom, with an asterisk denoting $p < 0.10$ after rounding to two decimals (p -values that round to 0.00 are shown as <0.01). No controls are included. Standard errors are clustered at the stratification level (107 clusters). The final row regresses the complier indicator on all 20 covariates and reports the cluster-robust F test that their coefficients are jointly zero; four covariates are exact linear combinations of the others (each set of education, income, left-right, and outlet dummies is complete), so the test covers the 16 identified coefficients, $F(16, 734)$.

Table A.18: Complier balance tests – outcome covariates

Family	Outcome	$\Delta Complier$	Std. Error	Test p -val (= 0)
News diet	1. News consumption (minutes)	0.07	0.09	0.4
	2. Closely follows politics	0.01	0.08	0.95
	3. Main medium: App/Sites	0.09	0.08	0.25
	4. Main medium: TV	-0.18	0.09	0.04*
	5. Main medium: Social media	0.07	0.08	0.37
	6. Intensity: Web news	0.23	0.08	<0.01*
	7. Intensity: Social media news	0.10	0.08	0.23
	8. Intensity: Cable TV news	-0.04	0.07	0.58
	9. Intensity: Local TV news	-0.03	0.07	0.71
	10. Intensity: National TV news	-0.02	0.08	0.83
	11. Intensity: Video news	-0.11	0.08	0.17
	12. Intensity: Radio news	-0.09	0.08	0.25
	13. Index: news diet	0.02	0.04	0.53
News knowledge	14. News knowledge	0.29	0.07	<0.01*
	15. Trust in chosen newspaper	-0.03	0.07	0.62
	16. Confidence in press	0.15	0.08	0.07*
Opinions about press	17. Fair media coverage of elections	0.01	0.09	0.88
	18. Fair media coverage of politics	0.02	0.09	0.84
Political engagement	19. Index: press opinions	0.07	0.08	0.36
	20. Talk about politics	0.05	0.08	0.49
	21. Share political content on WhatsApp	-0.06	0.09	0.48
	22. Share political content on social media	0.06	0.08	0.43
	23. Volunteers for campaigns	-0.00	0.09	0.99
	24. Donates to campaigns	0.01	0.07	0.94
	25. Tries to convince others to vote	-0.00	0.08	0.96
	26. Watches governor debates	-0.03	0.08	0.68
	27. Reads interviews with candidates	0.25	0.09	<0.01*
Willingness-to-vote	28. Index: political engagement	0.04	0.04	0.33
	29. Accepts election outcome	-0.19	0.09	0.05*
	30. Rule by executive orders	-0.20	0.08	0.01*
	31. Ban protests	-0.12	0.09	0.18
	32. Prosecute journalists	-0.06	0.09	0.48
Trust in electoral system	33. Ignore court rulings	0.02	0.07	0.81
	34. Index: willingness-to-vote	-0.11	0.07	0.12
	35. Elections are free and fair	0.16	0.10	0.1
	36. Trust in electronic voting system	0.22	0.09	0.01*
	37. Confidence in elections	0.26	0.09	<0.01*
Polarization	38. Index: trust in electoral system	0.20	0.08	0.01*
	39. Bolsonaro x Lula	0.05	0.09	0.59
	40. Bolsonaro x Supreme Court	0.08	0.10	0.4
	41. Left x Right politicians	0.08	0.08	0.36
	42. Left x Right supporters	-0.01	0.10	0.88
	Joint significance test			<0.01

Note: Differences in mean baseline outcomes between participants who redeemed their coupon (*compliers*) and those who did not, within the 751 participants assigned to either treatment arm (*S* and *SI* pooled; 526 compliers). The comparison is compliers against non-compliers, not *SI* against control. The outcomes are the 42 of Table A.16, and they are not on a common scale: among the 751, the items have standard deviations between 0.81 and 1.10, four of the family indices between 0.56 and 0.88, and the news-diet index (row 13) is unstandardized, with a standard deviation of 0.44. Column 1 reports the difference, from a regression of the outcome on a complier indicator; column 2 its standard error; and column 3 the p -value, from a t distribution with 749 degrees of freedom, with an asterisk denoting $p < 0.10$ after rounding to two decimals (p -values that round to 0.00 are shown as <0.01). No controls are included. Standard errors are clustered at the stratification level (107 clusters). The final row regresses the complier indicator on all 42 outcomes and reports the cluster-robust F test of the 38 identified coefficients, $F(38, 712)$: the press, engagement, willingness-to-vote, and trust indices (rows 19, 28, 34, and 38) are exact linear combinations of their component rows and are dropped.

Table A.19: Multiple hypothesis corrections – pre-registered list of outcomes

Outcome	p -val	BH q -val	BY q -val	dBH q -val
1. News consumption (minutes)	0.88	0.96	1	0.93
2. Index: news diet	<0.01	0.01	0.03	0.01
3. Index: press opinions	0.02	0.12	0.34	0.11
4. Interested in Black Friday offer	0.05	0.14	0.42	0.14
5. News knowledge	0.37	0.56	1	0.56
6. Sep. 7 news knowledge	0.2	0.34	1	0.34
7. Elections news knowledge	0.03	0.11	0.34	0.11
8. Political knowledge	0.85	1	1	0.93
9. Polarization Lula x Bolsonaro	0.92	0.92	1	0.93
10. Index: political engagement	0.11	0.25	0.79	0.26
11. Index: trust in electoral system	0.82	1	1	0.93
12. Index: willingness-to-vote	0.13	0.26	0.79	0.26

Note: One row per pre-registered outcome. Column 1 reports the unadjusted p -value on $\hat{\beta}_{SI}$ from Equation (1); columns 2–4 report q -values under the Benjamini–Hochberg (Benjamini and Hochberg, 1995), Benjamini–Yekutieli (Benjamini and Yekutieli, 2001), and dependence-adjusted (Fithian and Lei, 2022) corrections. A q -value is the smallest expected false discovery rate at which the null is rejected. No standard errors are reported; see Appendix E.

Table A.20: Multiple hypothesis corrections – main outcomes

Outcome	p -val	BH q -val	BY q -val	dBH q -val
1. News consumption (minutes)	0.88	1	1	0.98
2. Closely follows politics	0.04	0.23	0.9	0.21
3. Main medium: App/Sites	0.21	0.42	1	0.42
4. Main medium: TV	0.56	0.75	1	0.77
5. Main medium: Social media	0.05	0.21	0.9	0.21
6. Intensity: Web news	<0.01	<0.01	<0.01	0.01
7. Intensity: Social media news	0.09	0.25	1	0.26
8. Index: News diet	<0.01	0.01	0.06	0.02
9. Sep. 7 news knowledge	0.2	0.43	1	0.42
10. Elections news knowledge	0.03	0.2	0.85	0.19
11. News knowledge	0.37	0.7	1	0.67
12. Political knowledge	0.85	1	1	0.99
13. Index: information	0.06	0.22	0.98	0.23
14. Trust in chosen newspaper	<0.01	<0.01	<0.01	0.01
15. Confidence in press	0.36	0.7	1	0.66
16. Fair media coverage of elections	0.05	0.22	0.9	0.21
17. Fair media coverage of politics	0.02	0.19	0.74	0.16
18. Index: press opinions	0.02	0.17	0.74	0.17
19. Talk about politics	0.55	0.77	1	0.75
20. Share political content on WhatsApp	0.46	0.71	1	0.69
21. Share political content on social media	0.42	0.68	1	0.68
22. Volunteers for campaigns	0.42	0.7	1	0.68
23. Donates to campaigns	0.69	0.87	1	0.88
24. Tries to convince others to vote	0.09	0.27	1	0.25
25. Watches governor debates	0.09	0.29	1	0.25
26. Reads interviews with candidates	0.14	0.32	1	0.33
27. Index: political engagement	0.11	0.29	1	0.28
28. Accepts election outcome	0.48	0.68	1	0.7
29. Rule by executive orders	0.05	0.21	0.9	0.21
30. Ban protests	0.41	0.71	1	0.68
31. Prosecute journalists	0.96	0.96	1	0.98
32. Ignore court rulings	0.15	0.33	1	0.35
33. Temporary military intervention	0.41	0.73	1	0.67
34. Index: willingness-to-vote	0.13	0.32	1	0.32
35. Elections are free and fair	0.96	0.98	1	0.96
36. Trust in electronic voting system	0.96	1	1	0.99
37. Confidence in elections	0.47	0.7	1	0.68
38. Index: trust in electoral system	0.82	1	1	0.96
39. Polarization: Bolsonaro x Lula	0.92	0.99	1	0.97
40. Polarization: Bolsonaro x Supreme Court	0.68	0.89	1	0.85
41. Polarization: Left x Right politicians	0.91	1	1	0.98
42. Polarization: Left x Right supporters	0.88	1	1	0.98
43. Interested in Black Friday offer	0.05	0.24	0.9	0.21

Note: One row per outcome reported in the main text. Column 1 reports the unadjusted p -value on $\hat{\beta}_{SI}$ from Equation (1) and columns 2–4 the corresponding q -values under the three corrections. No standard errors are reported. See the note to Table A.19 for the definition of the corrections and Appendix E for details.

Table A.21: Heterogeneous treatment effects – above-median news consumption

	Web news intensity	Trust in newspaper	Press opinions index	Knowledge index	Affective polarization: Bolsonaro x Lula	Political engagement index	Trust in electoral system index	Willingness-to-vote index
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Subscription + Incentives (<i>SI</i>)	0.258 (0.058) [<0.01]	0.326 (0.055) [<0.01]	0.099 (0.046) [0.03]	0.151 (0.066) [0.02]	0.017 (0.051) [0.75]	0.074 (0.049) [0.13]	0.034 (0.040) [0.40]	−0.085 (0.053) [0.11]
Subscription + Incentives (<i>SI</i>) × Above Median News Consumption	−0.166 (0.103) [0.11]	−0.075 (0.100) [0.45]	−0.034 (0.083) [0.68]	−0.140 (0.113) [0.22]	−0.080 (0.086) [0.35]	−0.041 (0.088) [0.64]	−0.085 (0.067) [0.20]	0.064 (0.088) [0.47]
Control raw mean	1.81	2.42	0	0	70.04	0	0	0
Control raw std. dev.	1.22	0.98	1	1.02	32.51	0.99	1.01	1.02
Observations	1,363	1,363	1,363	1,088	1,363	1,363	1,363	1,363

Note: Estimates are $\hat{\beta}_2$ from Equation (2), the differential treatment effect between participants above and below the median in baseline news consumption. Standard errors are heteroskedasticity-robust. See the note to Table A.5 for conventions.

Table A.22: Heterogeneous treatment effects – above-median social media time

	Web news intensity	Trust in newspaper	Press opinions index	Knowledge index	Affective polarization: Bolsonaro x Lula	Political engagement index	Trust in electoral system index	Willingness-to-vote index
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Subscription + Incentives (<i>SI</i>)	0.240 (0.067) [<0.01]	0.325 (0.062) [<0.01]	0.143 (0.055) [<0.01]	0.185 (0.078) [0.02]	0.005 (0.060) [0.93]	0.086 (0.059) [0.15]	0.003 (0.045) [0.95]	−0.091 (0.061) [0.14]
Subscription + Incentives (<i>SI</i>) × Above Median Social Media Time	−0.073 (0.095) [0.45]	−0.075 (0.092) [0.42]	−0.108 (0.076) [0.15]	−0.193 (0.107) [0.07]	−0.009 (0.083) [0.92]	−0.048 (0.082) [0.56]	0.019 (0.063) [0.76]	0.051 (0.085) [0.55]
Control raw mean	1.81	2.42	0	0	70.04	0	0	0
Control raw std. dev.	1.22	0.98	1	1.02	32.51	0.99	1.01	1.02
Observations	1,363	1,363	1,363	1,088	1,363	1,363	1,363	1,363

Note: Estimates are $\hat{\beta}_2$ from Equation (2), the differential treatment effect between participants above and below the median in baseline time spent on social media. Standard errors are heteroskedasticity-robust. See the note to Table A.5 for conventions.

Table A.23: Heterogeneous treatment effects – above-median social media consumption

	Web news intensity	Trust in newspaper	Press opinions index	Knowledge index	Affective polarization: Bolsonaro x Lula	Political engagement index	Trust in electoral system index	Willingness-to-vote index
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Subscription + Incentives (<i>SI</i>)	0.196 (0.055) [<0.01]	0.304 (0.052) [<0.01]	0.087 (0.045) [0.05]	0.069 (0.062) [0.27]	−0.012 (0.048) [0.80]	0.042 (0.045) [0.35]	0.031 (0.037) [0.41]	−0.065 (0.049) [0.19]
Subscription + Incentives (<i>SI</i>) × Above Median Social Media Consumption	0.055 (0.114) [0.63]	−0.007 (0.111) [0.95]	0.020 (0.086) [0.81]	0.143 (0.133) [0.28]	0.039 (0.098) [0.69]	0.064 (0.102) [0.53]	−0.097 (0.071) [0.17]	−0.013 (0.098) [0.90]
Control raw mean	1.81	2.42	0	0	70.04	0	0	0
Control raw std. dev.	1.22	0.98	1	1.02	32.51	0.99	1.01	1.02
Observations	1,363	1,363	1,363	1,088	1,363	1,363	1,363	1,363

Note: Estimates are $\hat{\beta}_2$ from Equation (2), the differential treatment effect between participants above and below the median in baseline social media news consumption. Standard errors are heteroskedasticity-robust. See the note to Table A.5 for conventions.

Table A.24: Heterogeneous treatment effects – above-median news knowledge

	Web news intensity	Trust in newspaper	Press opinions index	Knowledge index	Affective polarization: Bolsonaro x Lula	Political engagement index	Trust in electoral system index	Willingness-to-vote index
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Subscription + Incentives (<i>SI</i>)	0.216 (0.064) [<0.01]	0.349 (0.062) [<0.01]	0.152 (0.052) [<0.01]	0.056 (0.070) [0.42]	−0.009 (0.062) [0.88]	0.016 (0.054) [0.76]	0.041 (0.043) [0.35]	0.011 (0.060) [0.86]
Subscription + Incentives (<i>SI</i>) × Above Median News Knowledge	−0.025 (0.097) [0.80]	−0.119 (0.092) [0.20]	−0.147 (0.077) [0.05]	0.109 (0.110) [0.32]	0.028 (0.081) [0.73]	0.100 (0.080) [0.21]	−0.100 (0.064) [0.12]	−0.165 (0.084) [0.05]
Control raw mean	1.81	2.42	0	0	70.04	0	0	0
Control raw std. dev.	1.22	0.98	1	1.02	32.51	0.99	1.01	1.02
Observations	1,363	1,363	1,363	1,088	1,363	1,363	1,363	1,363

Note: Estimates are $\hat{\beta}_2$ from Equation (2), the differential treatment effect between participants above and below the median in baseline news knowledge. Standard errors are heteroskedasticity-robust. See the note to Table A.5 for conventions.

Table A.25: Heterogeneous treatment effects – news choice

	Web news intensity	Trust in newspaper	Press opinions index	Knowledge index	Affective polarization: Bolsonaro x Lula	Political engagement index	Trust in electoral system index	Willingness-to-vote index
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Subscription + Incentives (<i>SI</i>) × Carta Capital	0.265 (0.097) [<0.01]	0.268 (0.087) [<0.01]	−0.001 (0.074) [0.99]	0.191 (0.123) [0.12]	−0.057 (0.062) [0.36]	0.219 (0.084) [<0.01]	−0.014 (0.061) [0.82]	−0.095 (0.086) [0.27]
Subscription + Incentives (<i>SI</i>) × Estadão	0.335 (0.094) [<0.01]	0.461 (0.088) [<0.01]	0.138 (0.077) [0.07]	−0.018 (0.108) [0.87]	0.047 (0.086) [0.59]	−0.039 (0.074) [0.60]	−0.010 (0.069) [0.89]	0.001 (0.087) [0.99]
Subscription + Incentives (<i>SI</i>) × Gazeta do Povo	0.464 (0.148) [<0.01]	0.496 (0.170) [<0.01]	0.224 (0.117) [0.06]	0.240 (0.172) [0.16]	0.254 (0.143) [0.08]	−0.112 (0.136) [0.41]	−0.144 (0.117) [0.22]	−0.210 (0.145) [0.15]
Subscription + Incentives (<i>SI</i>) × O Globo	−0.058 (0.094) [0.54]	0.149 (0.090) [0.10]	0.042 (0.079) [0.60]	−0.005 (0.102) [0.96]	−0.094 (0.085) [0.27]	0.018 (0.077) [0.82]	0.118 (0.058) [0.04]	−0.084 (0.084) [0.32]
Subscription + Incentives (<i>SI</i>) × Valor Econômico	0.195 (0.111) [0.08]	0.095 (0.106) [0.37]	0.120 (0.107) [0.26]	0.292 (0.162) [0.07]	−0.012 (0.115) [0.92]	0.214 (0.114) [0.06]	0.035 (0.095) [0.72]	0.035 (0.116) [0.76]
Control raw mean	1.81	2.42	0	0	70.04	0	0	0
Control raw std. dev.	1.22	0.98	1	1.02	32.51	0.99	1.01	1.02
Observations	1,363	1,363	1,363	1,088	1,363	1,363	1,363	1,363

Note: Estimates are $\hat{\beta}_h$ from Equation (3), the treatment effect within each outlet chosen at baseline; there is no omitted reference outlet, so the rows are levels and are not comparable pairwise. Standard errors are heteroskedasticity-robust. See the note to Table A.5 for conventions.

Table A.26: Heterogeneous treatment effects – political position

	Web news intensity	Trust in newspaper	Press opinions index	Knowledge index	Affective polarization: Bolsonaro x Lula	Political engagement index	Trust in electoral system index	Willingness-to-vote index
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Subscription + Incentives (SI) $\times$ Left	0.180 (0.065) [<0.01]	0.200 (0.056) [<0.01]	0.029 (0.047) [0.54]	0.077 (0.072) [0.29]	0.015 (0.043) [0.73]	0.126 (0.057) [0.03]	−0.043 (0.034) [0.20]	0.0002 (0.056) [1.00]
Subscription + Incentives (SI) $\times$ Center	0.297 (0.085) [<0.01]	0.248 (0.088) [<0.01]	0.176 (0.079) [0.03]	0.105 (0.099) [0.29]	0.050 (0.095) [0.60]	−0.018 (0.067) [0.79]	0.122 (0.071) [0.09]	−0.147 (0.085) [0.09]
Subscription + Incentives (SI) $\times$ Right	0.165 (0.124) [0.18]	0.771 (0.132) [<0.01]	0.178 (0.103) [0.08]	0.233 (0.128) [0.07]	−0.266 (0.130) [0.04]	0.050 (0.102) [0.62]	0.090 (0.090) [0.32]	−0.123 (0.106) [0.25]
Control raw mean	1.81	2.42	0	0	70.04	0	0	0
Control raw std. dev.	1.22	0.98	1	1.02	32.51	0.99	1.01	1.02
Observations	1,363	1,363	1,363	1,088	1,363	1,363	1,363	1,363

Note: Estimates are $\hat{\beta}_h$ from Equation (3), the treatment effect within each self-declared political position at baseline. Standard errors are heteroskedasticity-robust. See the note to Table A.5 for conventions.

Table A.27: Heterogeneous treatment effects – non-left versus left

	Web news intensity	Trust in newspaper	Press opinions index	Knowledge index	Affective polarization: Bolsonaro x Lula	Political engagement index	Trust in electoral system index	Willingness-to-vote index
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Subscription + Incentives (<i>SI</i>)	0.180 (0.065) [<0.01]	0.200 (0.056) [<0.01]	0.029 (0.046) [0.54]	0.077 (0.071) [0.28]	0.015 (0.043) [0.73]	0.126 (0.057) [0.03]	−0.043 (0.034) [0.20]	0.0002 (0.056) [1.00]
Subscription + Incentives (<i>SI</i>) × Non-left (center or right)	0.057 (0.095) [0.55]	0.227 (0.093) [0.01]	0.158 (0.078) [0.04]	0.063 (0.107) [0.55]	−0.067 (0.088) [0.44]	−0.147 (0.080) [0.07]	0.140 (0.067) [0.04]	−0.150 (0.086) [0.08]
Control raw mean	1.81	2.42	0	0	70.04	0	0	0
Control raw std. dev.	1.22	0.98	1	1.02	32.51	0.99	1.01	1.02
Observations	1,363	1,363	1,363	1,088	1,363	1,363	1,363	1,363

Note: Estimates are from Equation (2) with the moderator equal to one for participants who declared a centrist or right-wing position at baseline. The first row is the treatment effect among left-wingers; the second row is the non-left minus left differential, $\hat{\beta}_2$, whose p -values Section 6 reports. Standard errors are heteroskedasticity-robust. See the note to Table A.5 for conventions.

Table A.28: Heterogeneous treatment effects – above-median age

	Web news intensity	Trust in newspaper	Press opinions index	Knowledge index	Affective polarization: Bolsonaro x Lula	Political engagement index	Trust in electoral system index	Willingness-to-vote index
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Subscription + Incentives (<i>SI</i>)	0.290 (0.067) [<0.01]	0.306 (0.066) [<0.01]	0.115 (0.052) [0.03]	0.160 (0.075) [0.03]	0.060 (0.057) [0.29]	0.122 (0.059) [0.04]	0.054 (0.048) [0.26]	−0.108 (0.060) [0.07]
Subscription + Incentives (<i>SI</i>) × Above Median Age	−0.181 (0.095) [0.06]	−0.018 (0.092) [0.85]	−0.059 (0.076) [0.44]	−0.128 (0.109) [0.24]	−0.132 (0.083) [0.11]	−0.124 (0.081) [0.12]	−0.105 (0.064) [0.10]	0.104 (0.085) [0.22]
Control raw mean	1.81	2.42	0	0	70.04	0	0	0
Control raw std. dev.	1.22	0.98	1	1.02	32.51	0.99	1.01	1.02
Observations	1,363	1,363	1,363	1,088	1,363	1,363	1,363	1,363

Note: Estimates are $\hat{\beta}_2$ from Equation (2), the differential treatment effect between participants above and below the median age. Standard errors are heteroskedasticity-robust. See the note to Table A.5 for conventions.

Table A.29: Heterogeneous treatment effects – gender

	Web news intensity	Trust in newspaper	Press opinions index	Knowledge index	Affective polarization: Bolsonaro x Lula	Political engagement index	Trust in electoral system index	Willingness-to-vote index
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Subscription + Incentives (<i>SI</i>)	0.270 (0.056) [<0.01]	0.312 (0.053) [<0.01]	0.090 (0.045) [0.05]	0.124 (0.062) [0.05]	0.023 (0.050) [0.65]	0.076 (0.048) [0.12]	0.014 (0.038) [0.72]	−0.034 (0.051) [0.50]
Subscription + Incentives (<i>SI</i>) × Male (ref. Female)	−0.256 (0.110) [0.02]	−0.068 (0.104) [0.51]	−0.017 (0.085) [0.84]	−0.090 (0.123) [0.47]	−0.089 (0.088) [0.31]	−0.071 (0.087) [0.41]	−0.036 (0.066) [0.59]	−0.087 (0.094) [0.35]
Control raw mean	1.81	2.42	0	0	70.04	0	0	0
Control raw std. dev.	1.22	0.98	1	1.02	32.51	0.99	1.01	1.02
Observations	1,363	1,363	1,363	1,088	1,363	1,363	1,363	1,363

Note: Estimates are $\hat{\beta}_2$ from Equation (2), the differential treatment effect between male and female participants. Standard errors are heteroskedasticity-robust. See the note to Table A.5 for conventions.

Table A.30: Heterogeneous treatment effects – above-median household income

	Web news intensity	Trust in newspaper	Press opinions index	Knowledge index	Affective polarization: Bolsonaro x Lula	Political engagement index	Trust in electoral system index	Willingness-to-vote index
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Subscription + Incentives (<i>SI</i>)	0.274 (0.064) [<0.01]	0.355 (0.063) [<0.01]	0.117 (0.052) [0.02]	0.094 (0.072) [0.19]	0.012 (0.057) [0.83]	0.032 (0.052) [0.54]	−0.004 (0.044) [0.92]	−0.092 (0.060) [0.13]
Subscription + Incentives (<i>SI</i>) × Above Median Income	−0.174 (0.097) [0.07]	−0.124 (0.092) [0.18]	−0.065 (0.077) [0.40]	−0.007 (0.109) [0.95]	−0.027 (0.083) [0.74]	0.066 (0.081) [0.41]	0.030 (0.064) [0.64]	0.070 (0.085) [0.41]
Control raw mean	1.81	2.42	0	0	70.04	0	0	0
Control raw std. dev.	1.22	0.98	1	1.02	32.51	0.99	1.01	1.02
Observations	1,363	1,363	1,363	1,088	1,363	1,363	1,363	1,363

Note: Estimates are $\hat{\beta}_2$ from Equation (2), the differential treatment effect between participants above and below the median household income. Standard errors are heteroskedasticity-robust. See the note to Table A.5 for conventions.

B Additional experimental details

We describe the “user journey” throughout the experiment. Figure 2 presents the flowchart of the experiment, from the pre-treatment period, consisting of the screening and baseline surveys, to the treatment period up until the endline survey.

Screening survey. First, by clicking on the advertisement, the respondent opens a screening survey. In this survey, we check the eligibility criteria cited above. The criteria were assessed sequentially.

Once participants met all the criteria and agreed to participate, we proceeded with a brief explanation of the study tasks, which we group into three different tasks: i) the possible coupon redemption and account creation, in case they are chosen to receive one; ii) the two main surveys, a baseline survey to be distributed in two days and an endline survey to follow in three months; and iii) the phone surveys that will happen regularly through WhatsApp. We then present the consent form, obtain consent from the participants, and collect contact information (email, telephone number) and demographic information (income, race, sex, and education).

After the end of the screening survey, the 2,055 participants who passed screening were sent a “welcome to the study” message through WhatsApp to confirm that their telephone number was reachable.

Baseline survey. Two days after the screening survey, we send the baseline survey to the individual by email and WhatsApp. Participants had about a week to reply, and we sent two reminders over this period. Of the 2,055 screened participants, 1,658 completed the baseline survey, and these 1,658 constitute the baseline analysis frame. Twenty-three of them share an identical IP address together with identical exact GPS coordinates with another respondent, a pattern consistent with duplicate responses from a single device or household. We retain these respondents in the sample.

By the end of the baseline survey, we randomize the treatment: Control participants received an end-of-survey message. In contrast, those in treatment receive an extra section displaying their coupon and providing instructions on redeeming it for each newspaper.

Treatment implementation. In addition to the survey instructions, we also send an email and WhatsApp message containing a document with instructions for coupon redemption and account creation on the newspaper websites. One day after first sending the baseline survey message, we follow up with a reminder. A few days later, we also reminded those individuals who have not yet redeemed their coupon as per administrative data from the partner news outlets.

Phone surveys. We run three types of phone surveys: i) News Consumption surveys, ii) Incentives treatment quizzes, and iii) Event Knowledge quizzes. The News Consumption phone survey is answered by all participants and repeats questions about news consumption quantity and media (TV, radio, etc.). As mentioned earlier in the design, the incentives treatment quiz is only sent to the Subscription+Incentives (*SI*) group. It also happens biweekly and alternates with the news consumption survey.

Finally, the event knowledge quiz is triggered at the researchers’ discretion. Its objective

is to measure participants' awareness of ongoing political events being covered extensively by the media in the country. Possible examples include the arrest of a former minister due to corruption, the opening of an inquiry to investigate one of the candidates, or the start of a scandal due to the leakage of some private information.

Endline surveys. After three months, all participants receive the endline survey via email and WhatsApp. We also send two reminders, one and two days after the first message, to try to reduce attrition.

C Data appendix: variable definitions and survey wordings

Table A.31: Variable definitions and survey wording

Variable Name	Survey Questions	Measure
News consumption of chosen outlet (administrative data)		Google Analytics records provide a timestamp for each URL accessed on partner newspapers' websites or apps. A "browsing session" groups URL entries whenever consecutive entries are no more than 10 minutes apart. Minutes of use are measured by the timestamp difference between consecutive entries within a session. The records cover browsers and apps on computers and smartphones. We set a ceiling of 10 minutes on entry durations and impute the final entry of each session, whose duration is not observed. Control participants are assigned zero minutes of consumption.
News consumption (self-reported)		Self-reported minutes per day of news consumption, ranging from 0 to 120 minutes per day (winsorized at 120 minutes).
Qualitative measure of amount of news consumption	Thinking back over the last four weeks, how closely did you follow news about Brazil's politics?	Four-point scale from "Not at all closely" to "very closely."

News diet index	Over the past four weeks, how often did you watch/read/listen/get news from...?	Responses use a five-point scale from “Never” to “Very often”. The eight media are: cable TV news, local TV news, print newspapers, national evening network TV news, social media websites (examples given are Facebook, Instagram, and Twitter), radio/podcasts, websites or apps other than social media, and video platforms (examples given are TikTok, Kwai, and YouTube). We combine standardized responses into a signed inverse-covariance-weighted index. Social media and video enter with reversed signs, so the index increases when consumption shifts toward professional outlets.
Individual news diet	Over the past four weeks, how often did you watch/read/listen/get news from...?	Each of the eight individual components of the news diet index.
Confidence in press index	Confidence in the press is measured on a four-point scale from “Not at all” to “A lot.”	Self-reported confidence in the press, using the same four-point scale from “Not at all” to “A lot.”
Fair coverage of elections index	Fair coverage questions are measured on a four-point scale from “Not at all often” to “Very often”.	Self-reported assessment of whether journalists cover elections fairly, on the same four-point scale from “Not at all often” to “Very often”.
Fair coverage of politics index	Fair coverage questions are measured on a four-point scale from “Not at all often” to “Very often”.	Self-reported assessment of whether journalists cover politics fairly, on the same four-point scale from “Not at all often” to “Very often”.

“Positive opinions about press” index	Index of the three items above (confidence in the press, fair coverage of elections, fair coverage of politics); see their wording.	Index combining confidence in the press with assessments of whether journalists provide fair coverage of elections and politics.
Trust in the chosen news outlet	We ask respondents whether they trust the news reported by each of the partner newspapers.	Four-point scale from “None at all” to “A great deal”, plus “Unfamiliar”.
Perceived ideological position of chosen news outlet	We also ask respondents to classify the newspapers into ideological positions	Seven-point scale from “Extreme left” to “Extreme right”. Baseline and endline responses therefore measure changes in perceived ideology.
Willingness to continue subscribing	At the end of the experiment, we offered participants in both the control and treatment groups the option to renew or sign up for a subscription to retain access to the newspapers’ content.	Differential take-up between treatment and control measures “willingness to continue subscribing.” Prices vary across outlets and were not set by the authors. Because endline coincided with Black Friday promotions, offered prices were below average subscription prices.

News knowledge	Score on a ten-statement news knowledge quiz. Four were “true news,” three were “incorrect news,” and three were “fake news.”	For each of the ten statements, participants choose “True,” “False,” or “I don’t know.” Scores assign one point to a correct answer, zero points to an incorrect answer, and half a point to “I don’t know.” Statement types are: i) “True news”: relevant items reported in the press at most two weeks before the quiz; ii) “Incorrect news”: “true news” items modified slightly to become factually incorrect; iii) “Fake news”: actual, “malicious” fake news circulating at the time and selected for relevance to political events or the campaign (in partnership with a co-editor of <i>Estado Verifica</i>).
News knowledge: true news	Score on a news knowledge quiz. Four were “true news”	Same score as the “News knowledge” variable
News knowledge: incorrect news	Score on a news knowledge quiz. Three were “incorrect news”	Same score as the “News knowledge” variable
News knowledge: fake news	Score on a news knowledge quiz. Three were “fake news”	Same score as the “News knowledge” variable
Event knowledge and awareness	As major news events occurred during the study period, we asked whether participants had been exposed to and were aware of them. These measures are similar to the “News Knowledge” quizzes above but focus on specific events.	Same score as the “News knowledge” variable

Political Knowledge	The statements asked about the presidency of the Senate; the presidency of the Chamber of Deputies; the number of representatives per state in the Senate; the number of representatives per state in the Chamber of Deputies; the duration of a senator's mandate in Brazil; the number of votes needed to approve a constitutional amendment; the parties with the most representatives in the Chamber of Deputies; and the number of Supreme Court Justices appointed by incumbent President Jair Bolsonaro.	Participants choose "True," "False," or "I don't know" for each of the eight statements. Scoring assigns one point to a correct answer, zero points to an incorrect answer, and half a point to "I don't know."
Polarization	Feeling Thermometer	The "feeling thermometer" measures affect toward President Bolsonaro, former President Lula, right-wing supporters, left-wing supporters, right-wing politicians, left-wing politicians, and the Supreme Court. Polarization between Bolsonaro and Lula and between right- and left-wing supporters captures distance between the two main camps of Brazilian politics. We also measure polarization between left- and right-wing politicians and between the Supreme Court and Bolsonaro.

Political engagement index	Intensity of participation in the following political activities: talking about politics with friends and family, sharing political content on WhatsApp/Telegram, sharing political content on other social media, helping a candidate by volunteering for the campaign, helping a candidate by trying to persuade others to vote, helping a candidate by donating money to the campaign, watching debates for governor and president, and watching interviews with candidates.	Seven outcomes use a five-point frequency scale from “Never” to “Very often”; the campaign-donation item uses a four-point scale.
Trust in electoral system index: free and fair elections	We asked individuals their opinions regarding free and fair elections: “In your view, are federal elections held in Brazil free and fair?” Possible answers are “Yes” or “No.”	We standardize responses into an index measuring deviations from the control group.
Trust in electoral system index: trust in the electronic system	We asked individuals about their trust in the electronic voting system, using a four-point scale from “Do not trust” to “Trust completely.”	We standardize responses into an index measuring deviations from the control group.
Trust in electoral system index: trust in elections	We asked individuals about their confidence in elections, using a four-point scale from “None at all” to “A great deal.”	We standardize responses into an index measuring deviations from the control group.
Trust in electoral system index	We group the previous three indices into one index summarizing respondents’ trust in the electoral system in general	We standardize the combined responses as an index of deviations from the control group.

Willingness-to-vote for un-democratic politicians index	We presented statements about actions by a respondent's candidate and used a five-point scale from "much more willing" to "much less willing" to elicit willingness to vote for the candidate. The statements include: contesting electoral results if losing by a small margin; overusing executive orders if Congress is not cooperating; prosecuting journalists who reveal misconduct without revealing sources; not fulfilling judicial orders if the politician believes the judge was partial; not authorizing opposition protests in public spaces; and requesting military interventions to investigate claims of fraud or corruption.	The statements are combined into an index, as with the other outcomes. Each component of the "willingness-to-vote" preference index is also analyzed separately.
--	--	---

C.1 Browsing data: sources, harmonization, and quality

Partner-newspaper browsing records measure consumption, but they are not fully harmonized across outlets. Internal systems collect and process the Analytics data differently. Carta Capital used a soft paywall for a substantial share of its content, so users did not always need to log in. O Globo and Valor Econômico aggregate unique URLs within a browsing session (a roughly 30-minute window) and report browsing time directly from analytics; as a result, a session may contain only one homepage entry. Estadão, Gazeta do Povo, and Carta Capital instead provide records closer to Google Chrome histories, with one row per webpage access and no session-level grouping.

Browsing time is measured with error because Analytics infers it from differences between URL timestamps. Inactivity, standalone sessions—common when users arrive from social media—and behaviors such as opening several tabs quickly or reading through email can all distort duration. App records do not identify the specific URLs viewed within the app. Administrative consumption levels are therefore noisy measures.

Missing records are a third limitation. Outlet systems associate consumption with 526 coupons, but the browsing files contain only 322 unique coupons. Three newspapers record between 82% and 91% of redeemed coupons, consistent with some subscribers redeeming access without subsequently using the website or app. Coverage is much lower for Carta Capital (39%, or 45 out of 116 consumed subscriptions) and O Globo (35%, or 50 out of 142). The soft paywall may explain Carta Capital, but O Globo remains far below the roughly 86% average of the other three outlets despite similar consumption in other measures, including self-reports. Some O Globo participants may therefore be absent from the browsing data.²⁸

We audited collection quality by browsing about 100 unique URLs from each outlet on a computer over one week and asking the newspapers to extract those records. Chrome browsing histories provide the benchmark for the outlet extracts. Appendix Table A.3 reports the results. Every outlet misses some URLs, primarily homepage visits or category landing pages such as “World” and “Economy.”²⁹ Missing unique URLs are the most informative metric because they are less sensitive to the researchers’ browsing behavior. The missing share ranges from 0% (Estadão) to 11% (Carta Capital), averaging 5.5%. Weighting by treatment-group shares yields 6.1% missing unique URLs overall.

Aside from the two low-coverage newspapers, the audit indicates that outlet systems capture a high share of unique URLs but often omit reloads and some homepage visits. We therefore treat total browsing time as a lower-bound measure.

²⁸When contacted regarding the missing participants, the newspapers assured us that they extracted all data available, and we asked for new extractions to make sure this was indeed the case. Thus, there is no reason to believe there are human mistakes in the extractions.

²⁹Carta Capital does not record homepage accesses, so its number of missing URLs is considerably higher.

D Differential attrition

In our experiment, treatment-control response rates diverge starting with the first post-baseline survey. Table A.15, columns 1, 3, and 5, reports completion rates for the three post-baseline samples. The gap is stable: both *treatment* groups are between 12 and 15 pp. less likely to respond, equivalent to between 13% and 20% of the control group’s completion rate. This is unusual, as we usually expect the control group, who did not receive a free subscription, to be more subject to attrition. Therefore, being randomly assigned to receive a free news subscription induced non-response. Attrition is a concern because it can introduce bias in the estimation of treatment effects.

We assess differential attrition using a series of balance tests. First, conditional on survey completion, we show treatment and control respondents are balanced on baseline variables. We extend our balance analysis following Ghanem et al. (2026). Second, we show how key observables, such as our main outcomes, do not predict attrition and differential attrition. Third, we show a strong predictor of attrition is treatment take-up — redeeming the subscription coupon. This is especially problematic if take-up captures latent demand for news consumption. At baseline, we asked participants to estimate how their news consumption would change after receiving the treatment. We find the difference in expected treatment-induced consumption between compliers and non-compliers is not statistically distinguishable from zero, and the estimated gap is small relative to the baseline mean.

Differences in other baseline outcomes and demographics are at most 0.29 standard deviations, while observables scarcely predict compliance. We therefore view sign-up and coupon-redemption frictions as the more plausible explanation for early exit among non-compliers.

In the subsections below, we expand on the evidence provided above and then discuss our bounding procedure for treatment effects.

D.1 The post-baseline samples are balanced on key covariates and outcomes.

We first repeat the overall balance analysis in Table A.1 for each of the relevant post-baseline samples. This test asks whether the respondent subsamples remain balanced on baseline variables.

Overall, all post-baseline samples remain balanced on covariates. We summarize the results in Appendix Figure A.2, adding the balance tests for the baseline sample as a benchmark. For every variable and sample, the magnitude of the differences between control and treatment averages are similar to those found in baseline sample. The F -test of joint orthogonality, obtained by regressing a treatment dummy on all covariates, yields a p -value 0.29 when performed at the baseline sample, while the p -value is 0.58, 0.45 and 0.60 when we consider the endline, Pre-elections, and September 7 survey samples.

However, as Ghanem et al. (2026) shows, the test above does not leverage all data. In particular, it only uses information from the respondent subsample. Ghanem et al. (2026) provide a

complete analysis of differential attrition that leverages information from the non-respondent subsample, and assesses balance in primary outcomes measured at baseline. Denote the response status dummy by R_i and treatment status by T_i . The object of interest in our study is the $ATE - R$, defined as the treatment effect on the respondent subpopulation,

$$ATE_R = E[Y_i|T_i = 1, R_i = 1] - E[Y_i|T_i = 0, R_i = 1],$$

where for our main analysis, T_i is the indicator for the *SI* arm. Ghanem et al. (2026) show how a lack of differential attrition, usually assessed by the “equal attrition rates” test, is not necessary nor sufficient to identify the ATE_R . Instead, the ATE_R is identified if treatment assignment and potential outcomes are independent conditional on response status. This condition implies a sharp testable implication consisting of an equality in average outcomes, measured at baseline, between treatment and control groups, conditional on response status.

Denote by Y_{i0} an outcome variable measured at baseline. The sharp testable implication corresponds to $E[Y_{i0}|T_i, R_i] = E[Y_{i0}|R_i]$. Note that this implies leveraging information from both the respondent ($R_i = 1$) and non-respondent ($R_i = 0$) subpopulations. For stratified designs, the test for the $ATE - R$ consists of the null hypothesis

$$H_0 : \pi_{10} = \pi_{11} = 0, \tag{D.4}$$

where π_{10} and π_{11} are obtained from the following regression equation,

$$Y_{i0} = \sum_{s=1}^S (\alpha^s + \beta_{01}^s R_i) \mathbb{I}\{S_i = s\} + \pi_{10} T_i + \pi_{11} T_i R_i + \varepsilon_i. \tag{D.5}$$

This amounts to jointly testing whether the treatment and control groups are balanced in both the attritors and respondents subsamples, while allowing for overall differences between respondents and attritors.

We gather a comprehensive set of baseline outcomes Y_{i0} presented in the main body of the paper and perform the test outlined above. To provide a benchmark, we perform the same test on the baseline respondents – i.e., the whole study sample — where by default $R_i = 1$ for all participants, and thus the test consists only of $\pi_{10} = 0$, which would correspond to a standard balance test.

Table A.16 presents, for 42 outcomes measured at baseline, the p -values of the test (D.4) and the corresponding Benjamini-Hochberg q -values, as suggested by Ghanem et al. (2026), computed within each sample across the 42 outcomes. Columns 1 and 2 list the families and outcomes, columns 3 and 4 the baseline-sample benchmark, and the remaining three pairs are the endline, pre-elections, and September 7 samples.

At the 10% level, we reject the null for 7 outcomes in the baseline sample, 3 in the pre-elections sample, and 3 in the September 7 sample, and 8 in the endline sample, against 4 expected rejections. At the 5% level, we reject 3 in the baseline, 3 in the endline, 2 in the pre-elections, and 1 in the September 7 sample, against 2 expected. No p -value is below 0.01 in any sample. No q -value is below 0.20 in any sample, so none of the 42 individual tests is

rejected at a false discovery rate of 20%.

D.2 Attrition and differential attrition are not predicted by covariates.

The previous section examined the internal validity question of whether balance remains in the respondent subsamples. We now investigate if covariates predict attrition. We perform a prediction exercise with a LASSO logit model using a 10-fold cross-validation. We include all demographic variables measured in Appendix Figure A.2 and Appendix Table A.1, along with other baseline measures such as usage of specific social media apps.

We perform two prediction exercises. First, we analyze whether the covariates are predictive of attrition by performing a LASSO logit regression on the following equation:

$$Attrition_i = \gamma X_i + \varepsilon_i, \quad (D.6)$$

and we then analyze whether X are predictive of attrition if we let the covariates vary with treatment dummies,

$$Attrition_i = \gamma X_i + \eta X_i S_i + \alpha X_i S I_i + \varepsilon_i. \quad (D.7)$$

Both exercises have low predictive fit. For Equation (D.6), the pseudo- R^2 (the in-sample deviance ratio of a cross-validated lasso) is 0.07, 0.04, and 0.07 for the September 7, pre-elections, and endline survey samples, respectively. For Equation (D.7), the corresponding R^2 's are 0.09, 0.08, and 0.11 for the September 7, pre-elections, and endline survey samples, respectively.

Taken together, the results show how both attrition and differential attrition are relatively unpredicted by covariates, with pseudo- R^2 's of at most 0.11. This suggests attrition was mostly driven by variables not related to news consumption habits, political preferences, or demographics. However, the take-up of treatment is still correlated with attrition, which we turn to next.

D.3 Take-up is only weakly predicted by covariates.

Appendix Table A.15, columns 2, 4, and 6 suggest that differential attrition in the post-baseline samples is driven by participants who did not redeem their coupons for a free subscription – i.e., non-compliers.³⁰ Even though post-baseline samples are balanced on observables, investigating what predicts take-up patterns is useful to assess the extent of differential attrition bias.

We consider balance between compliers and non-compliers, using the same covariates in

³⁰In the causal inference literature, compliers are usually defined as those who take-up the treatment due to treatment assignment. Since treatment take-up is zero in the absence of treatment by design (i.e., our sample only contains individuals who did not subscribe), the complier set is the same as the set of participants who take-up the treatment.

our analysis of overall balance above and restricting the sample to exclude Control group participants:

$$Y_i = \beta_0 + \beta_1 \text{Complier}_i + \varepsilon_i$$

The results are presented separately in Appendix Table A.17 (demographic and news choice variables) and Appendix Table A.18 (outcome variables measured at baseline). The joint significance tests reject that all coefficients are jointly equal to zero, with both p -values being lower than 0.01. However, for both sets of covariates, the magnitude of the differences is at most 0.29 standard deviations. We also observe five individual p -vals below 0.01: age, web news, news knowledge, reads interviews, and confidence in elections. We also repeat the prediction exercise for compliance as a way to gauge the relative importance of the differences between compliers and non-compliers. As is the case with attrition, the pseudo- R^2 is relatively low, at 0.06.³¹ This suggests that observed baseline covariates have limited predictive power for compliance.

In summary, there are significant correlations between baseline variables and compliance, but overall compliance is not easily predicted from covariates, and is not large enough to drive differences in balances.

D.4 Potential drivers of the take-up decision

A first-order concern is that take-up is driven by demand for the treatment. Participants eager to read the newspapers are more likely to redeem their vouchers. But we would then also expect demand for treatment to correlate with other variables, such as trust in the newspapers, which we find no evidence of. Nevertheless, we take this possibility seriously.

At baseline, participants reported how much they expected a free subscription to increase their news consumption, which we use as a demand proxy. Regressing this measure on compliance shows that compliers expect a 2.1 minute, or 5% ($p = 0.26$), larger increase than non-compliers. The difference is mild and not statistically distinguishable from zero. Nor do baseline reasons for not subscribing—including time, interest, exposure, and price—differ between compliers and non-compliers.

We speculate two factors could drive non-compliance. First, there is a mechanical difficulty in creating an account in the online newspapers. Appendix Table A.17 suggests that compliers were 0.29 standard deviations younger than non-compliers (p -val: <0.01). These differences are significant at the 5% level, and they point in the direction of younger participants being better able to redeem their coupons. Furthermore, we have information on whether participants used their cell phones to answer the baseline survey. Redeeming the coupon was significantly more cumbersome on a mobile device. We find that mobile users are 14pp. ($p < 0.01$) less likely to redeem their coupons.

Second, non-compliance may be related to a lack of trust in the study. Since we offer a free

³¹This cross-validated pseudo- R^2 is mildly sensitive to the random fold assignment; we report the value under the study's global seed (500), and it ranges from 0.06 to 0.07 across seeds.

subscription, and given the prevalence of online scams in Brazil, participants may have been suspicious of whether the study was legitimate when they received the free coupon. Thus, they may have decided to cut their participation. This would be in line with the differential attrition pattern starting already in the September 7 phone survey. We attempted to qualitatively investigate differential attrition by randomly selecting five participants who were non-respondents in post-baseline surveys. We were blocked by four out of these five participants without receiving a reply.

D.5 Bounding results

In this section, we provide standard bounding exercises to the always-respondents treatment effects combining methods from Lee (2009) and Semenova (2025). Here, we consider Y_i to be a (possibly non-continuous) outcome, and R_i is the response status. $R(1)$ and $R(0)$ denote an observation’s potential response status, if the observation is and is not treated, respectively.

The seminal Lee (2009) approach takes as the parameter of interest the average treatment effect on the always-respondents, i.e., $E[Y(1) - Y(0) | R(1) = 1, R(0) = 1]$. We assume that $R(1) = 1 \implies R(0) = 1$ – that is, an individual who responds if treated would also do so if they were untreated, implying there are no *response compliers* in our dataset, only *response defiers*.³²

The procedure described in Lee (2009) entails trimming the control group for the best and worst-scenarios based on quantiles. Let $p_0 = \frac{E[R=1|D=1]}{E[R=1|D=0]}$ be the ratio of the treated to the control response rate; the proportion of control respondents trimmed is $1 - p_0$. The usual quantile-trimming is valid only for continuous Y variables; since many of our variables are discrete-valued, we trim outcomes based on sorting. Let i denote the ranking of the Y -variable, breaking ties randomly. We then sort the data by the ranking and trim the control group from above or below up to the defined proportion, defined by $i^L := (1 - p_0)N$ and $i^H := p_0N$. The use of bounding techniques for discrete-valued outcomes still retain their intuition of representing worst-case scenarios (for the lower bound) and best-case scenarios (for the upper bound) and can be informative, but the bounds can be skewed and often not cover the 95% confidence intervals.

Semenova (2025) generalize Lee bounds to a continuous variables and to conditional monotonicity, which uses covariate information to provide more precise bounds. How much tighter they are relative to the basic Lee bounds depend on how predictive the covariates X are of both R and Y . As we have seen in Section D.2, the covariates are in general unresponsive of R , so the improvements should be mild at best. Because the conditional trimming step relies on conditional quantiles, we compute these bounds for the 15 outcomes with at least ten distinct values (the indices, quiz scores, consumption minutes, and affective thermometers) and report the basic bound alone for the other items.

The bounds are displayed together in Appendix Figure A.4 for ease of exposition. We gather the main outcomes presented in Section 5, including the summary indices, along with the most

³²The typical monotonicity assumption is that $R(0) = 1 \implies R(1) = 1$, that is, any participant who would select into the sample if untreated would also do so if treated. This conforms with a typical situation where the treated participants have higher response rates. In our case, however, the opposite holds.

relevant news diet outcomes from Section 5.1. We display the point estimates and standard errors obtained from Equation 1 along with the basic Lee bounds (gray rectangles) and, for the continuous outcomes, the generalized bounds (colored rectangles).

The trimmed share is 13.7% of control respondents for outcomes measured at endline, 15.5% for the September 7 quiz, 17.2% for the pre-election quiz, and 18.2% for the knowledge index, whose components are not all observed for every endline respondent; item non-response counts as attrition throughout. This directly explains the width of the intervals: discarding a seventh of one arm moves the trimmed mean substantially, so for a standardized outcome the identified set spans between 0.12 and 0.65 standard deviations (median 0.45) even where the point estimate is close to zero.

None of the summary indices has generalized bounds that exclude zero. The news-diet index is bounded by $[-0.09, 0.26]$, the news-knowledge index by $[-0.07, 0.23]$, the press-opinions index by $[-0.12, 0.16]$, the political-engagement index by $[-0.05, 0.28]$, the willingness-to-vote index by $[-0.14, 0.11]$, and the trust-in-electoral-system index by $[-0.12, 0.05]$. The generalized bounds for each of the three news quizzes also include zero.

E Multiple hypothesis adjustments

Our multiple-testing procedure controls the false discovery rate (FDR) at the 5% level; equivalently, it limits the expected share of false discoveries to 5% (Romano et al., 2010).³³ Variance-covariance indices summarize groups of outcomes within families in the main text (Anderson, 2008).

The standard FDR procedures are Benjamini–Hochberg (BH; Benjamini and Hochberg, 1995) and the more conservative Benjamini–Yekutieli variant (BY; Benjamini and Yekutieli, 2001). Both operate on the vector of p -values but make different assumptions about dependence across tests. BH assumes independence; BY permits arbitrary dependence at the cost of statistical power.

Because our outcomes are correlated, our preferred adjustment uses the dependence-adjusted Benjamini–Hochberg (dBH) procedure (Fithian and Lei, 2022). It calibrates p -value thresholds using the variance-covariance matrix of the estimated coefficients. To estimate that matrix, we stack observations Y_i across outcomes and define X_i as the covariates selected by post-double selection for each outcome. A block-diagonal matrix of the X_i ’s then enters a seemingly unrelated regression (SUR). Errors are clustered at the individual level, recovering each participant’s cross-outcome error correlation—the dependence used by dBH—without imposing strata-level clustering. We apply dBH to the resulting variance-covariance matrix.

The pre-analysis plan specified multiple-testing corrections for twelve primary outcomes, chosen as the most relevant outcomes within each family. We report the outcome list and the corresponding corrections below. A few variable definitions changed; Appendix F documents those changes in detail. The families and outcomes on which we perform the corrections are:

- (1) News consumption (minutes), self-reported
- (2) News diet index
- (3) Positive opinions about press index
- (4) Willingness to continue subscribing (interest in subscribing)
- (5) News knowledge (at endline survey)
- (6) News knowledge (regarding the September 7 rallies)³⁴
- (7) News knowledge (regarding the presidential elections)
- (8) Political knowledge
- (9) Polarization Lula–Bolsonaro

³³“False discoveries” are rejections of true null hypotheses, which in our setting amounts to rejecting a null treatment effect.

³⁴An earlier version of this outcome was listed as “Event knowledge and awareness (measured by phone surveys)”. Instead of jointly considering the two events we covered with phone surveys, namely the September 7 rallies and the presidential elections, by bundling them in an index, we add both separately.

- (10) Political engagement index
- (11) Trust in electoral system index
- (12) Willingness-to-vote for democratic politicians index

We also apply the correction to the full set of 43 main-text outcomes. This set covers opinions about the press, information, polarization, democratic attitudes, interest in the Black Friday offer, and the principal news-diet measures: news minutes, attention to politics, main medium of consumption, and intensity of social-media and newspaper use. We exclude several secondary news-diet intensity outcomes (TV, radio, and videos).

Table A.19 reports results for the pre-registered outcomes; Table A.20 reports the corresponding results for the main-text outcomes.

F Changes to pre-registration

This section documents the changes between the pre-analysis plan (PAP) registered in the AEA RCT registry (AEARCTR-0009940) and the analyses presented in the paper. Outcome-definition changes were submitted before the endline survey. We organize the changes by topic.

Changes to outcome definitions.

- *Manipulation check (browsing data).* We added the definition of a “browsing session” to the measurement of time spent reading news in the chosen outlet. This definition is important to avoid overestimating browsing time, in addition to winsorizing long-duration entries.
- *News consumption and opinions.* We changed the secondary outcome “trust in the five partner newspapers” to “trust in the chosen news outlet,” as the latter is more meaningful given our design and more closely reflects participants’ experiences.
- *News knowledge quiz.* The quiz was expanded from eight to ten statements. Accordingly, the composition was changed to four “true news,” three “incorrect news,” and three “fake news” statements. We also made explicit that the “event knowledge quizzes” follow the same scoring rules as the “news knowledge” quiz.
- *Political knowledge quiz.* We explicitly listed the items covered in this quiz, which were not fully specified in the original PAP.
- *Political engagement index.* We made the following changes to the components of the index: (i) added “actively talking about politics with friends and family”; (ii) removed “attaching stickers to their cars/motorcycles,” “distributing pamphlets for a candidate,” “attending rallies,” and “changing profile picture in social media to support a candidate”; (iii) changed “actively sharing content for a candidate in WhatsApp/Telegram” and “actively sharing content for a candidate on other social media” to “sharing political content” on these platforms. These changes reflect adjustments to make the index more relevant to the digital media environment in which the study was conducted.
- *Willingness-to-vote for undemocratic politicians index.* We added a new statement about “suggesting a temporary military intervention to investigate a fraud or corruption claim,” which was motivated by the salience of this issue during the 2022 Brazilian elections.

Changes to the exposition and analysis.

- *Trust in the electoral system.* In the PAP, trust in the electoral system was grouped under the “democratic preferences” family of outcomes. In the main text, we present trust in the electoral system as a separate family, given the particular relevance and salience of this topic during the elections. This is a change in exposition only and does not affect the statistical analysis.

- *Browsing data collection.* The PAP pre-registered a form of data collection through a browsing extension on Google Chrome. Due to changes in Google’s Manifest V3 that limited the ability of extensions to export data to outside servers, this approach was not feasible. Instead, we relied on administrative browsing data provided directly by the partner newspapers, as described in Section 4.
- *Event knowledge surveys.* Rather than jointly considering the two events covered with phone surveys (the September 7 rallies and the presidential elections) by bundling them in a single index, we present results for each event separately, as this provides more informative evidence on the timing and nature of the treatment effects on knowledge.
- *Multiple hypothesis testing.* We pre-registered twelve primary outcomes for multiple hypothesis corrections. The list was slightly modified to accommodate the changes in outcome definitions described above, and to separate the two event knowledge quizzes. We also conducted a broader multiple hypothesis correction on the full set of 43 main outcomes presented in the paper. Details of the correction procedures are provided in Appendix E.
- *Response scales.* The PAP described the political engagement items on a four-point scale and trust in the electronic voting system on a three-point scale. The survey used a five-point scale for seven engagement items, a four-point scale for campaign donations, and a four-point scale for trust in the electronic voting system.
- *Analyses not in the PAP.* The following analyses were not pre-registered: the main-medium-of-consumption indicators, the index of trust in other news sources, question-level effects in the news quizzes and the rescoring of “I don’t know” answers, the pooled-arms comparison, and the attrition analysis in Appendix D.

Other changes.

- We corrected minor typographical errors in the PAP, including a reference to “nine mediums” when only eight were listed, and a typo in a section title.
- We added the specification that browsing data would only be obtained for users who opted into creating their newspaper accounts.
- The pre-endline, post-baseline pre-analysis plan displayed 1,664 observations; 6 observations were test responses that were incorrectly added, bringing the final number of baseline respondents to 1,658.